

# Copyright Notice

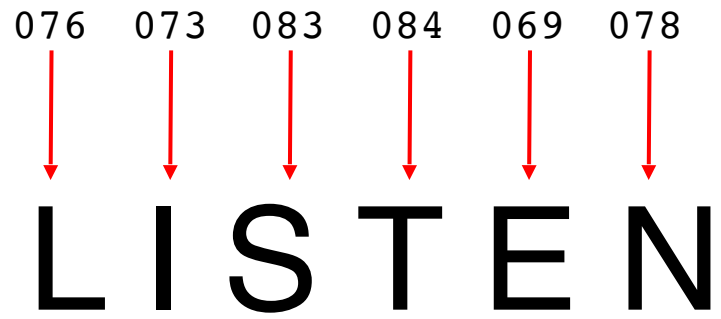
These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>

LISTEN

076 073 083 084 069 078



A diagram illustrating the word "LISTEN" with numerical values above each letter. The numbers are 076, 073, 083, 084, 069, and 078, corresponding to the letters L, I, S, T, E, and N respectively. Red arrows point from each number down to its corresponding letter.

Letter	Value
L	076
I	073
S	083
T	084
E	069
N	078

LISTEN

083 073 076 069 078 084  
↓ ↓ ↓ ↓ ↓ ↓  
S I L E N T

076 073 083 084 069 078  
↓ ↓ ↓ ↓ ↓ ↓  
L I S T E N

I love my dog

I love my dog



001

I love my dog

001



002



003



004



I love my dog

001

002

003

004

I love my cat



I love my dog

001

002

003

004

I love my cat

001

002

003

I love my dog

001

002

003

004

I love my cat

001

002

003

005



001

002

003

004

001

002

003

005

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.preprocessing.text import Tokenizer
```

```
sentences = [
    'I love my dog',
    'I love my cat'
]
```

```
tokenizer = Tokenizer(num_words = 100)
tokenizer.fit_on_texts(sentences)
word_index = tokenizer.word_index
print(word_index)
```

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.preprocessing.text import Tokenizer
```

```
sentences = [
    'I love my dog',
    'I love my cat'
]
```

```
tokenizer = Tokenizer(num_words = 100)
tokenizer.fit_on_texts(sentences)
word_index = tokenizer.word_index
print(word_index)
```

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.preprocessing.text import Tokenizer
```

```
sentences = [
    'I love my dog',
    'I love my cat'
]
```

```
tokenizer = Tokenizer(num_words = 100)
tokenizer.fit_on_texts(sentences)
word_index = tokenizer.word_index
print(word_index)
```

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.preprocessing.text import Tokenizer
```

```
sentences = [
    'I love my dog',
    'I love my cat'
]
```

```
tokenizer = Tokenizer(num_words = 100)
```

```
tokenizer.fit_on_texts(sentences)
```

```
word_index = tokenizer.word_index
```

```
print(word_index)
```

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.preprocessing.text import Tokenizer
```

```
sentences = [
    'I love my dog',
    'I love my cat'
]
```

```
tokenizer = Tokenizer(num_words = 100)
```

```
tokenizer.fit_on_texts(sentences)
```

```
word_index = tokenizer.word_index
```

```
print(word_index)
```



```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.preprocessing.text import Tokenizer
```

```
sentences = [
    'I love my dog',
    'I love my cat'
]
```

```
tokenizer = Tokenizer(num_words = 100)
tokenizer.fit_on_texts(sentences)
word_index = tokenizer.word_index
print(word_index)
```

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.preprocessing.text import Tokenizer
```

```
sentences = [
    'I love my dog',
    'I love my cat'
]
```

```
tokenizer = Tokenizer(num_words = 100)
tokenizer.fit_on_texts(sentences)
word_index = tokenizer.word_index
print(word_index)
```



```
sentences = [  
    'I love my dog',  
    'I love my cat',  
    'You love my dog!'  
]
```

```
sentences = [  
    'I love my dog',  
    'I love my cat',  
    'You love my dog!'  
]
```











```
from tensorflow.keras.preprocessing.text import Tokenizer
```

```
sentences = [  
    'I love my dog',  
    'I love my cat',  
    'You love my dog!',  
    'Do you think my dog is amazing?'  
]
```

```
tokenizer = Tokenizer(num_words = 100)
```

```
tokenizer.fit_on_texts(sentences)
```

```
word_index = tokenizer.word_index
```

```
sequences = tokenizer.texts_to_sequences(sentences)
```

```
print(word_index)
```

```
print(sequences)
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
```

```
sentences = [  
    'I love my dog',  
    'I love my cat',  
    'You love my dog!',  
    'Do you think my dog is amazing?'  
]
```

```
tokenizer = Tokenizer(num_words = 100)
```

```
tokenizer.fit_on_texts(sentences)
```

```
word_index = tokenizer.word_index
```

```
sequences = tokenizer.texts_to_sequences(sentences)
```

```
print(word_index)
```

```
print(sequences)
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
```

```
sentences = [  
    'I love my dog',  
    'I love my cat',  
    'You love my dog!',  
    'Do you think my dog is amazing?'  
]
```

```
tokenizer = Tokenizer(num_words = 100)  
tokenizer.fit_on_texts(sentences)  
word_index = tokenizer.word_index
```

```
sequences = tokenizer.texts_to_sequences(sentences)
```

```
print(word_index)  
print(sequences)
```

```
{'amazing': 10, 'dog': 3, 'you': 5, 'cat': 6,  
'think': 8, 'i': 4, 'is': 9, 'my': 1, 'do': 7,  
'love': 2}
```

```
[[4, 2, 1, 3], [4, 2, 1, 6], [5, 2, 1, 3], [7, 5, 8, 1, 3, 9, 10]]
```

```
{ 'amazing': 10, 'dog': 3, 'you': 5, 'cat': 6,  
  'think': 8, 'i': 4, 'is': 9, 'my': 1, 'do': 7,  
  'love': 2 }
```

```
[[4, 2, 1, 3], [4, 2, 1, 6], [5, 2, 1, 3], [7, 5, 8, 1, 3, 9, 10]]
```

```
{'amazing': 10, 'dog': 3, 'you': 5, 'cat': 6,  
'think': 8, 'i': 4, 'is': 9, 'my': 1, 'do': 7,  
'love': 2}
```

```
[[4, 2, 1, 3], [4, 2, 1, 6], [5, 2, 1, 3], [7, 5, 8, 1, 3, 9, 10]]
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
```

```
sentences = [  
    'I love my dog',  
    'I love my cat',  
    'You love my dog!',  
    'Do you think my dog is amazing?'  
]
```

```
tokenizer = Tokenizer(num_words = 100)
```

```
tokenizer.fit_on_texts(sentences)
```

```
word_index = tokenizer.word_index
```

```
sequences = tokenizer.texts_to_sequences(sentences)
```

```
print(word_index)
```

```
print(sequences)
```

```
test_data = [  
    'i really love my dog',  
    'my dog loves my manatee'  
]
```

```
test_seq = tokenizer.texts_to_sequences(test_data)  
print(test_seq)
```



```
test_data = [  
    'i really love my dog',  
    'my dog loves my manatee'  
]
```

```
test_seq = tokenizer.texts_to_sequences(test_data)  
print(test_seq)
```

```
[[4, 2, 1, 3], [1, 3, 1]]
```

```
{'think': 8, 'amazing': 10, 'my': 1, 'love': 2, 'dog': 3, 'is': 9, 'you': 5, 'do': 7, 'cat': 6, 'i': 4}
```

```
test_data = [  
    'i really love my dog',  
    'my dog loves my manatee'  
]  
  
test_seq = tokenizer.texts_to_sequences(test_data)  
print(test_seq)
```



```
[[4, 2, 1, 3], [1, 3, 1]]
```

```
{'think': 8, 'amazing': 10, 'my': 1, 'love': 2, 'dog': 3, 'is': 9, 'you': 5, 'do': 7, 'cat': 6, 'i': 4}
```

```
test_data = [  
    'i really love my dog',  
    'my dog loves my manatee'  
]  
  
test_seq = tokenizer.texts_to_sequences(test_data)  
print(test_seq)
```

```
[[4, 2, 1, 3], [1, 3, 1]]
```

```
{'think': 8, 'amazing': 10, 'my': 1, 'love': 2, 'dog': 3, 'is': 9, 'you': 5, 'do': 7, 'cat': 6, 'i': 4}
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
```

```
sentences = [  
    'I love my dog',  
    'I love my cat',  
    'You love my dog!',  
    'Do you think my dog is amazing?'  
]
```

```
tokenizer = Tokenizer(num_words = 100, oov_token="<OOV>")  
tokenizer.fit_on_texts(sentences)  
word_index = tokenizer.word_index
```

```
sequences = tokenizer.texts_to_sequences(sentences)
```

```
test_data = [  
    'i really love my dog',  
    'my dog loves my manatee'  
]
```

```
test_seq = tokenizer.texts_to_sequences(test_data)  
print(test_seq)
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
```

```
sentences = [  
    'I love my dog',  
    'I love my cat',  
    'You love my dog!',  
    'Do you think my dog is amazing?'  
]
```

```
tokenizer = Tokenizer(num_words = 100, oov_token="<OOV>")  
tokenizer.fit_on_texts(sentences)  
word_index = tokenizer.word_index
```



```
sequences = tokenizer.texts_to_sequences(sentences)
```

```
test_data = [  
    'i really love my dog',  
    'my dog loves my manatee'  
]
```

```
test_seq = tokenizer.texts_to_sequences(test_data)  
print(test_seq)
```

[[5, 1, 3, 2, 4], [2, 4, 1, 2, 1]]

{'think': 9, 'amazing': 11, 'dog': 4, 'do': 8, 'i': 5, 'cat': 7,  
'you': 6, 'love': 3, '<OOV>': 1, 'my': 2, 'is': 10}

```
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
sentences = [
    'I love my dog',
    'I love my cat',
    'You love my dog!',
    'Do you think my dog is amazing?'
]
```

```
tokenizer = Tokenizer(num_words = 100, oov_token="<OOV>")
tokenizer.fit_on_texts(sentences)
word_index = tokenizer.word_index
```

```
sequences = tokenizer.texts_to_sequences(sentences)
```

```
padded = pad_sequences(sequences)
print(word_index)
print(sequences)
print(padded)
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
sentences = [  
    'I love my dog',  
    'I love my cat',  
    'You love my dog!',  
    'Do you think my dog is amazing?'  
]
```

```
tokenizer = Tokenizer(num_words = 100, oov_token="<OOV>")  
tokenizer.fit_on_texts(sentences)  
word_index = tokenizer.word_index
```

```
sequences = tokenizer.texts_to_sequences(sentences)
```

```
padded = pad_sequences(sequences)  
print(word_index)  
print(sequences)  
print(padded)
```



```
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
sentences = [
    'I love my dog',
    'I love my cat',
    'You love my dog!',
    'Do you think my dog is amazing?'
]
```

```
tokenizer = Tokenizer(num_words = 100, oov_token="<OOV>")
tokenizer.fit_on_texts(sentences)
word_index = tokenizer.word_index
```

```
sequences = tokenizer.texts_to_sequences(sentences)
```

```
padded = pad_sequences(sequences)
```

```
print(word_index)
```

```
print(sequences)
```

```
print(padded)
```







```
padded = pad_sequences(sequences, padding='post')
```

```
padded = pad_sequences(sequences, padding='post', maxlen=5)
```




```
padded = pad_sequences(sequences, padding='post',  
                        truncating='post', maxlen=5)
```



Sarcasm in News Headlines Dataset by Rishabh Misra


<https://rishabhmisra.github.io/publications/>



 Dataset

## News Headlines Dataset For Sarcasm Detection

High quality dataset for the task of Sarcasm Detection


 Rishabh Misra · updated a year ago (Version 1)

154

^

[Data](#) [Kernels \(39\)](#) [Discussion \(2\)](#) [Activity](#) [Metadata](#) [Download \(2 MB\)](#) [New Kernel](#)

 CC0: Public Domain

 classification, deep learning, nlp, linguistics

### Description

#### Context

Past studies in Sarcasm Detection mostly make use of Twitter datasets collected using hashtag based supervision but such datasets are noisy in terms of labels and language. Furthermore, many tweets are replies to other tweets and detecting sarcasm in these requires the availability of contextual tweets.

To overcome the limitations related to noise in Twitter datasets, this **News Headlines dataset for Sarcasm Detection** is collected from two news website. [TheOnion](#) aims at producing sarcastic versions of current events and we collected all the headlines from News in Brief and News in Photos categories (which are sarcastic). We collect real (and non-sarcastic) news headlines from [HuffPost](#).

This new dataset has following advantages over the existing Twitter datasets:

- Since news headlines are written by professionals in a formal manner, there are no spelling mistakes and informal usage. This reduces the sparsity and also increases the chance of finding pre-trained embeddings.
- Furthermore, since the sole purpose of *TheOnion* is to publish sarcastic news, we get high-quality labels with much less noise as compared to Twitter datasets.
- Unlike tweets which are replies to other tweets, the news headlines we obtained are self-contained. This would help us in teasing apart the real sarcastic elements.

#### Content

Each record consists of three attributes:

- `is_sarcastic`: 1 if the record is sarcastic otherwise 0
- `headline`: the headline of the news article
- `article_link`: link to the original news article. Useful in collecting supplementary data

`is_sarcastic`: 1 if the record is sarcastic otherwise 0

`headline`: the headline of the news article

`article_link`: link to the original news article. Useful in collecting supplementary data

```
{"article_link": "https://politics.theonion.com/boehner-just-wants-wife-to-listen-not-come-up-with-alt-1819574302", "headline": "boehner just wants wife to listen, not come up with alternative debt-reduction ideas", "is_sarcastic": 1}
```

```
{"article_link": "https://www.huffingtonpost.com/entry/roseanne-revival-review_us_5ab3a497e4b054d118e04365", "headline": "the 'roseanne' revival catches up to our thorny political mood, for better and worse", "is_sarcastic": 0}
```

```
{"article_link": "https://local.theonion.com/mom-starting-to-fear-son-s-web-series-closest-thing-she-1819576697", "headline": "mom starting to fear son's web series closest thing she will have to grandchild", "is_sarcastic": 1}
```

[

```
{"article_link": "https://politics.theonion.com/boehner-just-wants-wife-to-listen-not-come-up-with-alt-1819574302", "headline": "boehner just wants wife to listen, not come up with alternative debt-reduction ideas", "is_sarcastic": 1},
```

```
{"article_link": "https://www.huffingtonpost.com/entry/roseanne-revival-review_us_5ab3a497e4b054d118e04365", "headline": "the 'roseanne' revival catches up to our thorny political mood, for better and worse", "is_sarcastic": 0},
```

```
{"article_link": "https://local.theonion.com/mom-starting-to-fear-son-s-web-series-closest-thing-she-1819576697", "headline": "mom starting to fear son's web series closest thing she will have to grandchild", "is_sarcastic": 1}
```

]

```
import json
```

```
with open("sarcasm.json", 'r') as f:  
    datastore = json.load(f)
```

```
sentences = []
```

```
labels = []
```

```
urls = []
```

```
for item in datastore:
```

```
    sentences.append(item['headline'])
```

```
    labels.append(item['is_sarcastic'])
```

```
    urls.append(item['article_link'])
```

```
import json
```

```
with open("sarcasm.json", 'r') as f:  
    datastore = json.load(f)
```

```
sentences = []
```

```
labels = []
```

```
urls = []
```

```
for item in datastore:
```

```
    sentences.append(item['headline'])
```

```
    labels.append(item['is_sarcastic'])
```

```
    urls.append(item['article_link'])
```

```
import json
```

```
with open("sarcasm.json", 'r') as f:  
    datastore = json.load(f)
```

```
sentences = []
```

```
labels = []
```

```
urls = []
```

```
for item in datastore:
```

```
    sentences.append(item['headline'])
```

```
    labels.append(item['is_sarcastic'])
```

```
    urls.append(item['article_link'])
```

```
import json
```

```
with open("sarcasm.json", 'r') as f:  
    datastore = json.load(f)
```

```
sentences = []
```

```
labels = []
```

```
urls = []
```

```
for item in datastore:
```

```
    sentences.append(item['headline'])
```

```
    labels.append(item['is_sarcastic'])
```

```
    urls.append(item['article_link'])
```

```
import json
```

```
with open("sarcasm.json", 'r') as f:  
    datastore = json.load(f)
```

```
sentences = []
```

```
labels = []
```

```
urls = []
```

```
for item in datastore:
```

```
    sentences.append(item['headline'])
```

```
    labels.append(item['is_sarcastic'])
```

```
    urls.append(item['article_link'])
```



```
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences

tokenizer = Tokenizer(oov_token="<OOV>")
tokenizer.fit_on_texts(sentences)
word_index = tokenizer.word_index

sequences = tokenizer.texts_to_sequences(sentences)
padded = pad_sequences(sequences, padding='post')
print(padded[0])
print(padded.shape)
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences

tokenizer = Tokenizer(oov_token="<OOV>")
tokenizer.fit_on_texts(sentences)
word_index = tokenizer.word_index
```

```
sequences = tokenizer.texts_to_sequences(sentences)
padded = pad_sequences(sequences, padding='post')
print(padded[0])
print(padded.shape)
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
tokenizer = Tokenizer(oov_token="<OOV>")
tokenizer.fit_on_texts(sentences)
word_index = tokenizer.word_index
```

```
sequences = tokenizer.texts_to_sequences(sentences)
padded = pad_sequences(sequences, padding='post')
print(padded[0])
print(padded.shape)
```

{'underwood': 24127, 'skillingsbolle': 23055, 'grabs': 12293, 'mobility': 8909, '"assassin's": 12648, 'visualize': 23973, 'hurting': 4992, 'orphaned': 9173, '"agreed"'': 24365, 'narration': 28470

```
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
tokenizer = Tokenizer(oov_token="<OOV>")
tokenizer.fit_on_texts(sentences)
word_index = tokenizer.word_index
```

```
sequences = tokenizer.texts_to_sequences(sentences)
padded = pad_sequences(sequences, padding='post')
print(padded[0])
print(padded.shape)
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences

tokenizer = Tokenizer(oov_token="<OOV>")
tokenizer.fit_on_texts(sentences)
word_index = tokenizer.word_index

sequences = tokenizer.texts_to_sequences(sentences)
padded = pad_sequences(sequences, padding='post')
print(padded[0])
print(padded.shape)
```

```
[ 308 15115  679 3337 2298  48  382 2576 15116  6 2577 8434
  0   0   0   0   0   0   0   0   0   0   0
  0   0   0   0   0   0   0   0   0   0   0
  0   0   0   0]
```

(26709, 40)

308	15115	679	3337	2298	48	382	2576	15116	6	2577	8434
0	0	0	0	0	0	0	0	0	0	0	
0	0	0	0	0	0	0	0	0	0	0	
0	0	0	0]								

(26709, 40)



```
[ 308 15115 679 3337 2298 48 382 2576 15116 6 2577 8434
  0 0 0 0 0 0 0 0 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0
  0 0 0 0]
```

(26709, 40)

# Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>

audio	"fashion_mnist"	text
	"horses_or_humans"	
"nsynth"	"image_label_folder"	"cnn_dailymail"
	"imagenet2012"	"glue"
image	"imagenet2012_corrupted"	"imdb_reviews"
	"kmnist"	"lm1b"
"abstract_reasoning"	"lsun"	"multi_nli"
"caltech101"	"mnist"	"squad"
"cats_vs_dogs"	"omniglot"	"wikipedia"
"celeb_a"	"open_images_v4"	"xnli"
"celeb_a_hq"	"oxford_iiit_pet"	
"cifar10"	"quickdraw_bitmap"	translate
"cifar100"	"rock_paper_scissors"	
"cifar10_corrupted"	"shapes3d"	"flores"
"coco2014"	"smallnorb"	"para_crawl"
"colorectal_histology"	"sun397"	"ted_hrlr_translate"
"cycle_gan"	"svhn_cropped"	"ted_multi_translate"
"diabetic_retinopathy..."	"tf_flowers"	"wmt15_translate"
"dsprites"		"wmt16_translate"
"dtd"	structured	"wmt17_translate"
"emnist"		"wmt18_translate"
	"higgs"	"wmt19_translate"
	"iris"	
	"titanic"	

audio	"fashion_mnist"
	"horses_or_humans"
"nsynth"	"image_label_folder"
	"imagenet2012"
image	"imagenet2012_corrupted"
	"kmnist"
"abstract_reasoning"	"lsun"
"caltech101"	"mnist"
"cats_vs_dogs"	"omniglot"
"celeb_a"	"open_images_v4"
"celeb_a_hq"	"oxford_iiit_pet"
"cifar10"	"quickdraw_bitmap"
"cifar100"	"rock_paper_scissors"
"cifar10_corrupted"	"shapes3d"
"coco2014"	"smallnorb"
"colorectal_histology"	"sun397"
"cycle_gan"	"svhn_cropped"
"diabetic_retinopathy..."	"tf_flowers"
"dsprites"	
"dtd"	structured
"emnist"	
	"higgs"
	"iris"
	"titanic"

text
"cnn_dailymail"
"glue"
"imdb_reviews"
"lm1b"
"multi_nli"
"squad"
"wikipedia"
"xnli"

translate
"flores"
"para_crawl"
"ted_hrlr_translate"
"ted_multi_translate"
"wmt15_translate"
"wmt16_translate"
"wmt17_translate"
"wmt18_translate"
"wmt19_translate"

<http://ai.stanford.edu/~amaas/data/sentiment/>

```
@InProceedings{maas-EtAl:2011:ACL-HLT2011,  
  author      = {Maas, Andrew L. and Daly, Raymond E. and Pham, Peter T. and Huang, Dan and  
Ng, Andrew Y. and Potts, Christopher},  
  title       = {Learning Word Vectors for Sentiment Analysis},  
  booktitle   = {Proceedings of the 49th Annual Meeting of the Association for Computational  
Linguistics: Human Language Technologies},  
  month       = {June},  
  year        = {2011},  
  address     = {Portland, Oregon, USA},  
  publisher   = {Association for Computational Linguistics},  
  pages       = {142--150},  
  url         = {http://www.aclweb.org/anthology/P11-1015}  
}
```

```
import tensorflow as tf  
print(tf.__version__)
```

```
import tensorflow_datasets as tfds
```

```
imdb, info = tfds.load("imdb_reviews", with_info=True, as_supervised=True)
```

```
import numpy as np
```

```
train_data, test_data = imdb['train'], imdb['test']
```



```
training_sentences = []
```

```
training_labels = []
```

```
testing_sentences = []
```

```
testing_labels = []
```

```
for s,l in train_data:
```

```
    training_sentences.append(str(s.numpy()))
```

```
    training_labels.append(l.numpy())
```

```
for s,l in test_data:
```

```
    testing_sentences.append(str(s.numpy()))
```

```
    testing_labels.append(l.numpy())
```

```
training_sentences = []
```

```
training_labels = []
```

```
testing_sentences = []
```

```
testing_labels = []
```

```
for s, l in train_data:
```

```
    training_sentences.append(str(s.numpy()))
```

```
    training_labels.append(l.numpy())
```

```
for s, l in test_data:
```

```
    testing_sentences.append(str(s.numpy()))
```

```
    testing_labels.append(l.numpy())
```

```
training_sentences = []
```

```
training_labels = []
```

```
testing_sentences = []
```

```
testing_labels = []
```

```
for s, l in train_data:
```

```
    training_sentences.append(str(s.numpy()))
```

```
    training_labels.append(l.numpy())
```

```
for s, l in test_data:
```

```
    testing_sentences.append(str(s.numpy()))
```

```
    testing_labels.append(l.numpy())
```

```
training_sentences = []
```

```
training_labels = []
```

```
testing_sentences = []
```

```
testing_labels = []
```

```
for s,l in train_data:
```

```
    training_sentences.append(str(s.numpy()))
```

```
    training_labels.append(l.numpy())
```

```
for s,l in test_data:
```

```
    testing_sentences.append(str(s.numpy()))
```

```
    testing_labels.append(l.numpy())
```

```
tf.Tensor(b"As a lifelong fan of Dickens, I have invariably been disappointed  
by adaptations of his novels.<br /><br />Although his works presented an  
extremely accurate re-telling of human life at every level in Victorian Britain,  
throughout them all was a pervasive thread of humour that could be both playful  
or sarcastic as the narrative dictated. In a way, he was a literary  
caricaturist and cartoonist. He could be serious and hilarious in the same  
sentence. He pricked pride, lampooned arrogance, celebrated modesty,  
and empathised with loneliness and poverty. It may be a cliché, but  
he was a people's writer.<br /><br />And it is the comedy that is so often  
missing from his interpretations. At the time of writing, Oliver Twist  
is being dramatised in serial form on BBC television. All of the misery  
and cruelty is their, but non of the humour, irony, and savage lampoonery.",  
shape=(), dtype=string)
```

```
tf.Tensor(1, shape=(), dtype=int64)
tf.Tensor(1, shape=(), dtype=int64)
tf.Tensor(1, shape=(), dtype=int64)
tf.Tensor(0, shape=(), dtype=int64)
tf.Tensor(0, shape=(), dtype=int64)
tf.Tensor(1, shape=(), dtype=int64)
```

```
training_labels_final = np.array(training_labels)
```

```
testing_labels_final = np.array(testing_labels)
```

```
vocab_size = 10000
embedding_dim = 16
max_length = 120
trunc_type='post'
oov_tok = "<OOV>"
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
tokenizer = Tokenizer(num_words = vocab_size, oov_token=oov_tok)
tokenizer.fit_on_texts(training_sentences)
word_index = tokenizer.word_index
sequences = tokenizer.texts_to_sequences(training_sentences)
padded = pad_sequences(sequences,maxlen=max_length, truncating=trunc_type)
```

```
testing_sequences = tokenizer.texts_to_sequences(testing_sentences)
testing_padded = pad_sequences(testing_sequences,maxlen=max_length)
```



```
vocab_size = 10000  
embedding_dim = 16  
max_length = 120  
trunc_type='post'  
oov_tok = "<OOV>"
```

```
from tensorflow.keras.preprocessing.text import Tokenizer  
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
tokenizer = Tokenizer(num_words = vocab_size, oov_token=oov_tok)  
tokenizer.fit_on_texts(training_sentences)  
word_index = tokenizer.word_index  
sequences = tokenizer.texts_to_sequences(training_sentences)  
padded = pad_sequences(sequences,maxlen=max_length, truncating=trunc_type)
```

```
testing_sequences = tokenizer.texts_to_sequences(testing_sentences)  
testing_padded = pad_sequences(testing_sequences,maxlen=max_length)
```

```
vocab_size = 10000
embedding_dim = 16
max_length = 120
trunc_type='post'
oov_tok = "<OOV>"
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
tokenizer = Tokenizer(num_words = vocab_size, oov_token=oov_tok)
tokenizer.fit_on_texts(training_sentences)
word_index = tokenizer.word_index
sequences = tokenizer.texts_to_sequences(training_sentences)
padded = pad_sequences(sequences,maxlen=max_length, truncating=trunc_type)

testing_sequences = tokenizer.texts_to_sequences(testing_sentences)
testing_padded = pad_sequences(testing_sequences,maxlen=max_length)
```

```
vocab_size = 10000
embedding_dim = 16
max_length = 120
trunc_type='post'
oov_tok = "<OOV>"
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
tokenizer = Tokenizer(num_words = vocab_size, oov_token=oov_tok)
```

```
tokenizer.fit_on_texts(training_sentences)
```

```
word_index = tokenizer.word_index
```

```
sequences = tokenizer.texts_to_sequences(training_sentences)
```

```
padded = pad_sequences(sequences,maxlen=max_length, truncating=trunc_type)
```

```
testing_sequences = tokenizer.texts_to_sequences(testing_sentences)
```

```
testing_padded = pad_sequences(testing_sequences,maxlen=max_length)
```

```
vocab_size = 10000
embedding_dim = 16
max_length = 120
trunc_type='post'
oov_tok = "<OOV>"
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
tokenizer = Tokenizer(num_words = vocab_size, oov_token=oov_tok)
```

```
tokenizer.fit_on_texts(training_sentences)
```

```
word_index = tokenizer.word_index
```

```
sequences = tokenizer.texts_to_sequences(training_sentences)
```

```
padded = pad_sequences(sequences,maxlen=max_length, truncating=trunc_type)
```

```
testing_sequences = tokenizer.texts_to_sequences(testing_sentences)
```

```
testing_padded = pad_sequences(testing_sequences,maxlen=max_length)
```

```
vocab_size = 10000
embedding_dim = 16
max_length = 120
trunc_type='post'
oov_tok = "<OOV>"
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
tokenizer = Tokenizer(num_words = vocab_size, oov_token=oov_tok)
tokenizer.fit_on_texts(training_sentences)
word_index = tokenizer.word_index
```

```
sequences = tokenizer.texts_to_sequences(training_sentences)
```

```
padded = pad_sequences(sequences,maxlen=max_length, truncating=trunc_type)
```

```
testing_sequences = tokenizer.texts_to_sequences(testing_sentences)
```

```
testing_padded = pad_sequences(testing_sequences,maxlen=max_length)
```

```
vocab_size = 10000
embedding_dim = 16
max_length = 120
trunc_type='post'
oov_tok = "<OOV>"
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
tokenizer = Tokenizer(num_words = vocab_size, oov_token=oov_tok)
tokenizer.fit_on_texts(training_sentences)
word_index = tokenizer.word_index
sequences = tokenizer.texts_to_sequences(training_sentences)
padded = pad_sequences(sequences,maxlen=max_length, truncating=trunc_type)
```

```
testing_sequences = tokenizer.texts_to_sequences(testing_sentences)
testing_padded = pad_sequences(testing_sequences,maxlen=max_length)
```

```
vocab_size = 10000
embedding_dim = 16
max_length = 120
trunc_type='post'
oov_tok = "<OOV>"
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
tokenizer = Tokenizer(num_words = vocab_size, oov_token=oov_tok)
tokenizer.fit_on_texts(training_sentences)
word_index = tokenizer.word_index
sequences = tokenizer.texts_to_sequences(training_sentences)
padded = pad_sequences(sequences,maxlen=max_length, truncating=trunc_type)
```

```
testing_sequences = tokenizer.texts_to_sequences(testing_sentences)
testing_padded = pad_sequences(testing_sequences,maxlen=max_length)
```

```
model = tf.keras.Sequential([  
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),  
    tf.keras.layers.Flatten(),  
    tf.keras.layers.Dense(6, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])
```



```
model = tf.keras.Sequential([  
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),  
    tf.keras.layers.Flatten(),  
    tf.keras.layers.Dense(6, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])
```

```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(6, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

```
model = tf.keras.Sequential([  
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),  
    tf.keras.layers.Flatten(),  
    tf.keras.layers.Dense(6, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])
```

Layer (type)	Output Shape	Param #
=====		
embedding_9 (Embedding)	(None, 120, 16)	160000
=====		
flatten_3 (Flatten)	(None, 1920)	0
=====		
dense_14 (Dense)	(None, 6)	11526
=====		
dense_15 (Dense)	(None, 1)	7
=====		
Total params: 171,533		
Trainable params: 171,533		
Non-trainable params: 0		

```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),
    tf.keras.layers.GlobalAveragePooling1D(),
    tf.keras.layers.Dense(6, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

Layer (type)	Output Shape	Param #
=====		
embedding_11 (Embedding)	(None, 120, 16)	160000
<hr/>		
global_average_pooling1d_3 ( (None, 16)	0	
<hr/>		
dense_16 (Dense)	(None, 6)	102
<hr/>		
dense_17 (Dense)	(None, 1)	7
=====		

Total params: 160,109  
Trainable params: 160,109  
Non-trainable params: 0

```
model.compile(loss='binary_crossentropy',optimizer='adam',metrics=['accuracy'])  
model.summary()
```

```
num_epochs = 10
model.fit(padded,
          training_labels_final,
          epochs=num_epochs,
          validation_data=(testing_padded, testing_labels_final))
```



Epoch 8/10

25000/25000 [=====] -

6s 256us/sample - loss: 5.2086e-04 - acc: 1.0000 - val\_loss: 0.7252 - val\_acc: 0.8270

Epoch 9/10

25000/25000 [=====] -

6s 222us/sample - loss: 3.0199e-04 - acc: 1.0000 - val\_loss: 0.7628 - val\_acc: 0.8269

Epoch 10/10

25000/25000 [=====] -

6s 224us/sample - loss: 1.7872e-04 - acc: 1.0000 - val\_loss: 0.7997 - val\_acc: 0.8259


```
e = model.layers[0]  
weights = e.get_weights()[0]  
print(weights.shape) # shape: (vocab_size, embedding_dim)
```

```
(10000, 16)
```

Hello : 1  
World : 2  
How : 3  
Are : 4  
You : 5

reverse\_word\_index = tokenizer.index\_word

1 : Hello  
2 : World  
3 : How  
4 : Are  
5 : You



```
import io
```

```
out_v = io.open('vecs.tsv', 'w', encoding='utf-8')
```

```
out_m = io.open('meta.tsv', 'w', encoding='utf-8')
```

```
for word_num in range(1, vocab_size):
```

```
    word = reverse_word_index[word_num]
```

```
    embeddings = weights[word_num]
```

```
    out_m.write(word + "\n")
```

```
    out_v.write('\t'.join([str(x) for x in embeddings]) + "\n")
```

```
out_v.close()
```

```
out_m.close()
```

```
import io
```

```
out_v = io.open('vecs.tsv', 'w', encoding='utf-8')
```

```
out_m = io.open('meta.tsv', 'w', encoding='utf-8')
```

```
for word_num in range(1, vocab_size):
```

```
    word = reverse_word_index[word_num]
```

```
    embeddings = weights[word_num]
```

```
    out_m.write(word + "\n")
```

```
    out_v.write('\t'.join([str(x) for x in embeddings]) + "\n")
```

```
out_v.close()
```

```
out_m.close()
```

```
import io
```

```
out_v = io.open('vecs.tsv', 'w', encoding='utf-8')
```

```
out_m = io.open('meta.tsv', 'w', encoding='utf-8')
```

```
for word_num in range(1, vocab_size):
```

```
    word = reverse_word_index[word_num]
```

```
    embeddings = weights[word_num]
```

```
    out_m.write(word + "\n")
```

```
    out_v.write("\t".join([str(x) for x in embeddings]) + "\n")
```

```
out_v.close()
```

```
out_m.close()
```

```
try:
    from google.colab import files
except ImportError:
    pass
else:
    files.download('vecs.tsv')
    files.download('meta.tsv')
```

## Embedding Projector



### DATA

5 tensors found

Word2Vec 10K ▾

Label by ▾

word

Color by ▾

No color map

☒ Sphereize data ⓘ

Load data

Publish

Checkpoint: Demo datasets

Metadata: oss\_data/word2vec\_10000\_200d\_labels.tsv

T-SNE

**PCA**

CUSTOM

X  
Component #1 ▾

Y  
Component #2 ▾

Z  
Component #3 ▾



PCA is approximate. ⓘ

Total variance described: 8.5%.



Points: 10000 | Dimension: 200



Show All  
Data

Isolate  
selection

Clear  
selection

Search



by

word ▾

BOOKMARKS (0) ⓘ





← → ↺

https://projector.tensorflow.org

☆ 🟢 📄 🔄 📁 🌈 🛡️ 🔍 🖨️ 🗑️ ⌵

Embedding Projector

?

🐙

DATA

🖨️ 🌙 A | Points: 10000 | Dimension: 200

5 tensors found

Word2Vec 10K

Label by

word

Color by

No color map

☒ Sphereize data ⓘ

Load data

Publish

Checkpoint: Demo datasets

Metadata: oss\_data/word2vec\_10000\_200d\_labels.tsv

T-SNE

PCA

CUSTOM

X

Component #1

Y

Component #2

Z

Component #3

☒

PCA is approximate. ⓘ

Total variance described: 8.5%.

?

🏠

Show All Data

Isolate selection

Clear selection

Search

by

word

BOOKMARKS (0) ⓘ

⌵

## Load data from your computer

### Step 1: Load a TSV file of vectors.

Example of 3 vectors with dimension 4:

```
0.1\t0.2\t0.5\t0.9  
0.2\t0.1\t5.0\t0.2  
0.4\t0.1\t7.0\t0.8
```

Choose file

### Step 2 (optional): Load a TSV file of metadata.

Example of 3 data points and 2 columns.

*Note: If there is more than one column, the first row will be parsed as column labels.*

```
Pokémon\tSpecies  
Wartortle\tTurtle  
Venusaur\tSeed  
Charmeleon\tFlame
```

Choose file

Click outside to dismiss.

## DATA

5 tensors found

Word2Vec 10K

☐ Sphereize data ?

Load data

Publish

Checkpoint: vecs.tsv

Metadata: meta.tsv

T-SNE

PCA

CUSTOM


X  
Component #1Y  
Component #2Z  
Component #3

PCA is approximate. ?

Total variance described: 98.7%.



Points: 9999 | Dimension: 16

Show All  
DataIsolate  
selectionClear  
selection

Search



by



BOOKMARKS (0) ?



## DATA

5 tensors found

Word2Vec 10K

☒ Sphereize data ?

Load data

Publish

Checkpoint: vecs.tsv

Metadata: meta.tsv

T-SNE

PCA

CUSTOM

x

Component #1

y

Component #2

z

Component #3



PCA is approximate. ?

Total variance described: 88.7%.




Points: 9999 | Dimension: 16

Show All  
DataIsolate  
selectionClear  
selection

Search



by



BOOKMARKS (0) ?



```
import json
```

```
import tensorflow as tf
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
```

```
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
vocab_size = 10000  
embedding_dim = 16  
max_length = 32  
trunc_type='post'  
padding_type='post'  
oov_tok = "<OOV>"  
training_size = 20000
```

```
!wget --no-check-certificate \  
  https://storage.googleapis.com/laurencemoroney-blog.appspot.com/sarcasm.json \  
  -O /tmp/sarcasm.json
```

```
with open("/tmp/sarcasm.json", 'r') as f:  
    datastore = json.load(f)
```

```
sentences = []  
labels = []
```

```
for item in datastore:  
    sentences.append(item['headline'])  
    labels.append(item['is_sarcastic'])
```



```
training_sentences = sentences[0:training_size]  
testing_sentences = sentences[training_size:]  
training_labels = labels[0:training_size]  
testing_labels = labels[training_size:]
```

```
training_sentences = sentences[0:training_size]  
testing_sentences = sentences[training_size:]  
training_labels = labels[0:training_size]  
testing_labels = labels[training_size:]
```

```
training_sentences = sentences[0:training_size]  
testing_sentences = sentences[training_size:]  
training_labels = labels[0:training_size]  
testing_labels = labels[training_size:]
```

```
training_sentences = sentences[0:training_size]
testing_sentences = sentences[training_size:]
training_labels = labels[0:training_size]
testing_labels = labels[training_size:]
```

```
tokenizer = Tokenizer(num_words=vocab_size, oov_token=oov_tok)
```

```
tokenizer.fit_on_texts(training_sentences)
```

```
word_index = tokenizer.word_index
```

```
training_sequences = tokenizer.texts_to_sequences(training_sentences)
```

```
training_padded = pad_sequences(training_sequences, maxlen=max_length,  
                                padding=padding_type, truncating=trunc_type)
```

```
testing_sequences = tokenizer.texts_to_sequences(testing_sentences)
```

```
testing_padded = pad_sequences(testing_sequences, maxlen=max_length,  
                               padding=padding_type, truncating=trunc_type)
```

```
tokenizer = Tokenizer(num_words=vocab_size, oov_token=oov_tok)
```

```
tokenizer.fit_on_texts(training_sentences)
```

```
word_index = tokenizer.word_index
```

```
training_sequences = tokenizer.texts_to_sequences(training_sentences)
```

```
training_padded = pad_sequences(training_sequences, maxlen=max_length,  
                                padding=padding_type, truncating=trunc_type)
```

```
testing_sequences = tokenizer.texts_to_sequences(testing_sentences)
```

```
testing_padded = pad_sequences(testing_sequences, maxlen=max_length,  
                               padding=padding_type, truncating=trunc_type)
```

```
tokenizer = Tokenizer(num_words=vocab_size, oov_token=oov_tok)
```

```
tokenizer.fit_on_texts(training_sentences)
```

```
word_index = tokenizer.word_index
```

```
training_sequences = tokenizer.texts_to_sequences(training_sentences)
```

```
training_padded = pad_sequences(training_sequences, maxlen=max_length,  
                                padding=padding_type, truncating=trunc_type)
```

```
testing_sequences = tokenizer.texts_to_sequences(testing_sentences)
```

```
testing_padded = pad_sequences(testing_sequences, maxlen=max_length,  
                               padding=padding_type, truncating=trunc_type)
```

```
tokenizer = Tokenizer(num_words=vocab_size, oov_token=oov_tok)
```

```
tokenizer.fit_on_texts(training_sentences)
```

```
word_index = tokenizer.word_index
```

```
training_sequences = tokenizer.texts_to_sequences(training_sentences)
```

```
training_padded = pad_sequences(training_sequences, maxlen=max_length,  
                                padding=padding_type, truncating=trunc_type)
```

```
testing_sequences = tokenizer.texts_to_sequences(testing_sentences)
```

```
testing_padded = pad_sequences(testing_sequences, maxlen=max_length,  
                               padding=padding_type, truncating=trunc_type)
```



```
tokenizer = Tokenizer(num_words=vocab_size, oov_token=oov_tok)
```

```
tokenizer.fit_on_texts(training_sentences)
```

```
word_index = tokenizer.word_index
```

```
training_sequences = tokenizer.texts_to_sequences(training_sentences)
```

```
training_padded = pad_sequences(training_sequences, maxlen=max_length,  
                                padding=padding_type, truncating=trunc_type)
```

```
testing_sequences = tokenizer.texts_to_sequences(testing_sentences)
```

```
testing_padded = pad_sequences(testing_sequences, maxlen=max_length,  
                               padding=padding_type, truncating=trunc_type)
```

```
tokenizer = Tokenizer(num_words=vocab_size, oov_token=oov_tok)
```

```
tokenizer.fit_on_texts(training_sentences)
```

```
word_index = tokenizer.word_index
```

```
training_sequences = tokenizer.texts_to_sequences(training_sentences)
```

```
training_padded = pad_sequences(training_sequences, maxlen=max_length,  
                                padding=padding_type, truncating=trunc_type)
```

```
testing_sequences = tokenizer.texts_to_sequences(testing_sentences)
```

```
testing_padded = pad_sequences(testing_sequences, maxlen=max_length,  
                                padding=padding_type, truncating=trunc_type)
```

```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),
    tf.keras.layers.GlobalAveragePooling1D(),
    tf.keras.layers.Dense(24, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

model.summary()

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 32, 16)	160000
global_average_pooling1d_2 (GlobalAveragePooling1D)	(None, 16)	0
dense_4 (Dense)	(None, 24)	408
dense_5 (Dense)	(None, 1)	25

Total params: 160,433  
Trainable params: 160,433  
Non-trainable params: 0



```
num_epochs = 30
```

```
history = model.fit(training_padded, training_labels, epochs=num_epochs,  
                    validation_data=(testing_padded, testing_labels), verbose=2)
```

```
import matplotlib.pyplot as plt
```

```
def plot_graphs(history, string):  
    plt.plot(history.history[string])  
    plt.plot(history.history['val_'+string])  
    plt.xlabel("Epochs")  
    plt.ylabel(string)  
    plt.legend([string, 'val_'+string])  
    plt.show()
```

```
plot_graphs(history, "acc")  
plot_graphs(history, "loss")
```



vocab\_size = 1000 (was 10,000)

embedding\_dim = 16

max\_length = 16 (was 32)



trunc\_type='post'

padding\_type='post'

oov\_tok = "<OOV>"

training\_size = 20000





vocab\_size = 1000 (was 10,000)

embedding\_dim = 32 (was 16)



max\_length = 16 (was 32)

trunc\_type='post'

padding\_type='post'

oov\_tok = "<OOV>"

training\_size = 20000



<https://github.com/tensorflow/datasets/tree/master/docs/catalog>

275 lines (209 sloc) | 9.65 KB

<> [icon] Raw Blame [icon] [icon] [icon]

## imdb\_reviews

- Description:

Large Movie Review Dataset. This is a dataset for binary sentiment classification containing substantially more data than previous benchmark datasets. We provide a set of 25,000 highly polar movie reviews for training, and 25,000 for testing. There is additional unlabeled data for use as well.

- Homepage: <http://ai.stanford.edu/~amaas/data/sentiment/>

- Source code: `tfds.text.IMDBReviews`

- Versions:

- `1.0.0` (default): New split API (<https://tensorflow.org/datasets/splits>)

- Download size: `80.23 MiB`

- Dataset size: `Unknown size`

- Auto-cached ([documentation](#)): Unknown

- Splits:

Split	Examples
<code>'test'</code>	25,000
<code>'train'</code>	25,000
<code>'unsupervised'</code>	50,000

275 lines (209 sloc) | 9.65 KB

<> [icon] Raw Blame [icon] [icon] [icon]

## imdb\_reviews/plain\_text (default config)

- Config description: Plain text
- Features:

```
FeaturesDict({  
  'label': ClassLabel(shape=(), dtype=tf.int64, num_classes=2),  
  'text': Text(shape=(), dtype=tf.string),  
})
```

- Examples ([tfds.as\\_dataframe](#)):

{% framebox %}

Display examples...

```
<script> const url = "https://storage.googleapis.com/tfds-data/visualization/dataframe/imdb_reviews-plain_text-1.0.0.html"; const dataButton  
= document.getElementById('displaydataframe'); dataButton.addEventListener('click', async () => { // Disable the button after clicking  
(dataframe loaded only once). dataButton.disabled = true;  
const contentPane = document.getElementById('dataframecontent'); try { const response = await fetch(url); // Error response codes don't throw  
an error, so force an error to show // the error message. if (!response.ok) throw Error(response.statusText);
```

```
const data = await response.text();  
contentPane.innerHTML = data;
```

```
} catch (e) { contentPane.innerHTML = 'Error loading examples. If the error persist, please open ' + 'a new issue.'; } }); </script>
```

{% endframebox %}

275 lines (209 sloc) | 9.65 KB

<> 📄 Raw Blame 🖨️ 📋 ✎️ 🗑️

## imdb\_reviews/bytes

- **Config description:** Uses byte-level text encoding with `tfds.deprecated.text.ByteTextEncoder`

- **Features:**

```
FeaturesDict({
  'label': ClassLabel(shape=(), dtype=tf.int64, num_classes=2),
  'text': Text(shape=(None,), dtype=tf.int64, encoder=<ByteTextEncoder vocab_size=257>),
})
```

- **Examples** ([tfds.as\\_dataframe](#)):

{% framebox %}

Display examples...

```
<script> const url = "https://storage.googleapis.com/tfds-data/visualization/dataframe/imdb_reviews-bytes-1.0.0.html"; const dataButton =
document.getElementById('displaydataframe'); dataButton.addEventListener('click', async () => { // Disable the button after clicking (dataframe
loaded only once). dataButton.disabled = true;
const contentPane = document.getElementById('dataframecontent'); try { const response = await fetch(url); // Error response codes don't throw
an error, so force an error to show // the error message. if (!response.ok) throw Error(response.statusText);
```

```
const data = await response.text();
contentPane.innerHTML = data;
```

```
} catch (e) { contentPane.innerHTML = 'Error loading examples. If the error persist, please open ' + 'a new issue.'; } }); </script>
```

{% endframebox %}

imdb\_reviews/bytes

275 lines (209 sloc) | 9.65 KB

<> [icon] Raw Blame [icon] [icon] [icon]

## imdb\_reviews/subwords8k

- **Config description:** Uses `tfds.deprecated.text.SubwordTextEncoder` with 8k vocab size
- **Features:**

```
FeaturesDict({  
  'label': ClassLabel(shape=(), dtype=tf.int64, num_classes=2),  
  'text': Text(shape=(None,), dtype=tf.int64, encoder=<SubwordTextEncoder vocab_size=8185>),  
})
```

- **Examples** ([tfds.as\\_dataframe](#)):

{% framebox %}

Display examples...

```
<script> const url = "https://storage.googleapis.com/tfds-data/visualization/dataframe/imdb_reviews-subwords8k-1.0.0.html"; const  
dataButton = document.getElementById('displaydataframe'); dataButton.addEventListener('click', async () => { // Disable the button after  
clicking (dataframe loaded only once). dataButton.disabled = true;  
const contentPane = document.getElementById('dataframecontent'); try { const response = await fetch(url); // Error response codes don't throw  
an error, so force an error to show // the error message. if (!response.ok) throw Error(response.statusText);
```

```
const data = await response.text();  
contentPane.innerHTML = data;
```

```
} catch (e) { contentPane.innerHTML = 'Error loading examples. If the error persist, please open ' + 'a new issue.'; } }); </script>
```

{% endframebox %}



```
import tensorflow_datasets as tfds
imdb, info = tfds.load("imdb_reviews/subwords8k", with_info=True, as_supervised=True)
```

```
train_data, test_data = imdb['train'], imdb['test']
```

```
tokenizer = info.features['text'].encoder
```

[tensorflow.org/datasets/api\\_docs/python/tfds/features/text/SubwordTextEncoder](https://tensorflow.org/datasets/api_docs/python/tfds/features/text/SubwordTextEncoder)

```
print(tokenizer_subwords.subwords)
```

```
['the_', ',', '.', 'a_', 'and_', 'of_', 'to_', 's_', 'is_', 'br', 'in_', 'l_', 'that_', 'this_', 'it_', ... ]
```

```
sample_string = 'TensorFlow, from basics to mastery'
```

```
tokenized_string = tokenizer_subwords.encode(sample_string)
```

```
print ('Tokenized string is {}'.format(tokenized_string))
```

```
original_string = tokenizer_subwords.decode(tokenized_string)
```

```
print ('The original string: {}'.format(original_string))
```

Tokenized string is [6307, 2327, 4043, 2120, 2, 48, 4249, 4429, 7, 2652, 8050]

The original string: TensorFlow, from basics to mastery

```
sample_string = 'TensorFlow, from basics to mastery'
```

```
tokenized_string = tokenizer_subwords.encode(sample_string)
```

```
print ('Tokenized string is {}'.format(tokenized_string))
```

```
original_string = tokenizer_subwords.decode(tokenized_string)
```

```
print ('The original string: {}'.format(original_string))
```

Tokenized string is [6307, 2327, 4043, 2120, 2, 48, 4249, 4429, 7, 2652, 8050]

The original string: TensorFlow, from basics to mastery

```
sample_string = 'TensorFlow, from basics to mastery'
```

```
tokenized_string = tokenizer_subwords.encode(sample_string)
```

```
print ('Tokenized string is {}'.format(tokenized_string))
```

```
original_string = tokenizer_subwords.decode(tokenized_string)
```

```
print ('The original string: {}'.format(original_string))
```

Tokenized string is [6307, 2327, 4043, 2120, 2, 48, 4249, 4429, 7, 2652, 8050]

The original string: TensorFlow, from basics to mastery

```
sample_string = 'TensorFlow, from basics to mastery'
```

```
tokenized_string = tokenizer_subwords.encode(sample_string)
```

```
print ('Tokenized string is {}'.format(tokenized_string))
```

```
original_string = tokenizer_subwords.decode(tokenized_string)
```

```
print ('The original string: {}'.format(original_string))
```

```
Tokenized string is [6307, 2327, 4043, 2120, 2, 48, 4249, 4429, 7, 2652, 8050]
```

```
The original string: TensorFlow, from basics to mastery
```



```
for ts in tokenized_string:  
    print ('{} ----> {}'.format(ts, tokenizer_subwords.decode([ts])))
```

6307 ----> Ten

2327 ----> sor

4043 ----> FI

2120 ----> ow

2 ----> ,

48 ----> from

4249 ----> basi

4429 ----> cs

7 ----> to

2652 ----> master

8050 ----> y

```
embedding_dim = 64
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(tokenizer_subwords.vocab_size, embedding_dim),
    tf.keras.layers.GlobalAveragePooling1D(),
    tf.keras.layers.Dense(6, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

model.summary()
```

```
embedding_dim = 64
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(tokenizer_subwords.vocab_size, embedding_dim),
    tf.keras.layers.GlobalAveragePooling1D(),
    tf.keras.layers.Dense(6, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

model.summary()
```

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, None, 64)	523840
global_average_pooling1d_1 (GlobalAveragePooling1D)	(None, 64)	0
dense_4 (Dense)	(None, 6)	390
dense_5 (Dense)	(None, 1)	7
Total params: 524,237		
Trainable params: 524,237		
Non-trainable params: 0		

```
num_epochs = 10
```

```
model.compile(loss='binary_crossentropy',  
              optimizer='adam',  
              metrics=['accuracy'])
```



```
history = model.fit(train_dataset,  
                    epochs=num_epochs,  
                    validation_data=test_data)
```

```
import matplotlib.pyplot as plt
```

```
def plot_graphs(history, string):  
    plt.plot(history.history[string])  
    plt.plot(history.history['val_'+string])  
    plt.xlabel("Epochs")  
    plt.ylabel(string)  
    plt.legend([string, 'val_'+string])  
    plt.show()
```

```
plot_graphs(history, "accuracy")
```

```
plot_graphs(history, "loss")
```




# Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>





$$f \left( \begin{array}{|c|} \hline \text{Data} \\ \hline \end{array} \begin{array}{|c|} \hline \text{Labels} \\ \hline \end{array} \right) = \text{Rules}$$

1

2

3

5

8

13

21

34

55

89

1

2

3

5

8

13

21

34

55

89

$n_0$

$n_1$

$n_2$

$n_3$

$n_4$

$n_5$

$n_6$

$n_7$


$n_8$


$n_9$


1	2	3	5	8	13	21	34	55	89
---	---	---	---	---	----	----	----	----	----

$n_0$	$n_1$	$n_2$	$n_3$	$n_4$	$n_5$	$n_6$	$n_7$	$n_8$	$n_9$
-------	-------	-------	-------	-------	-------	-------	-------	-------	-------


$$n_x = n_{x-1} + n_{x-2}$$

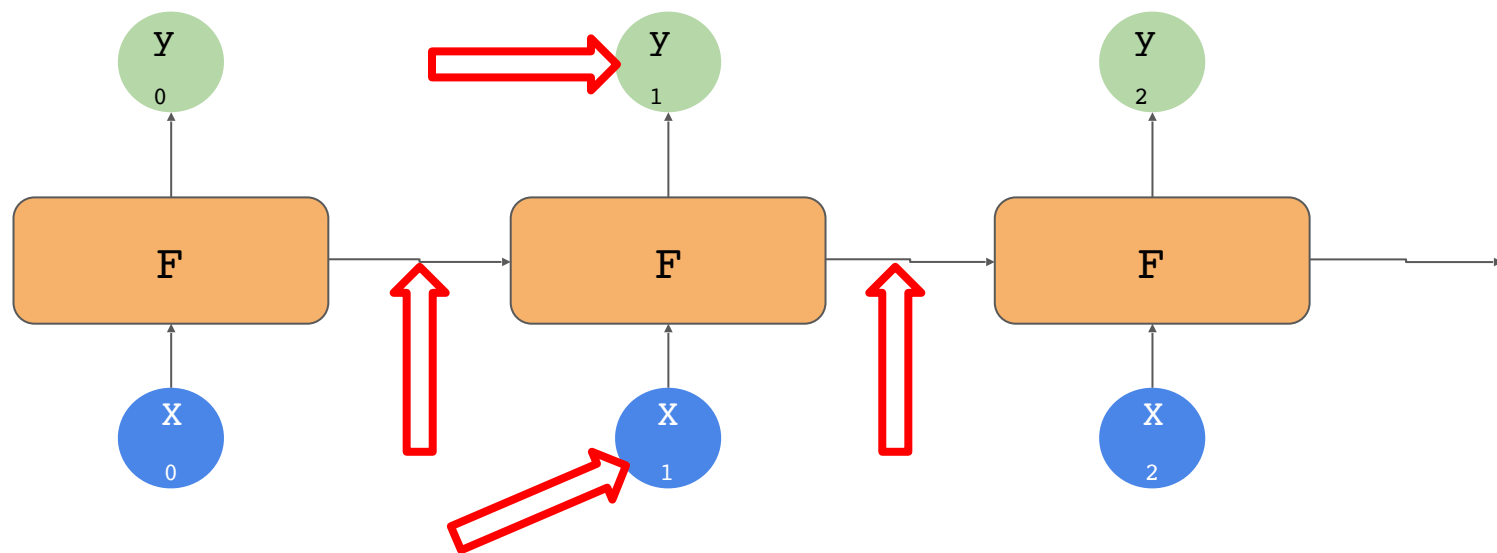


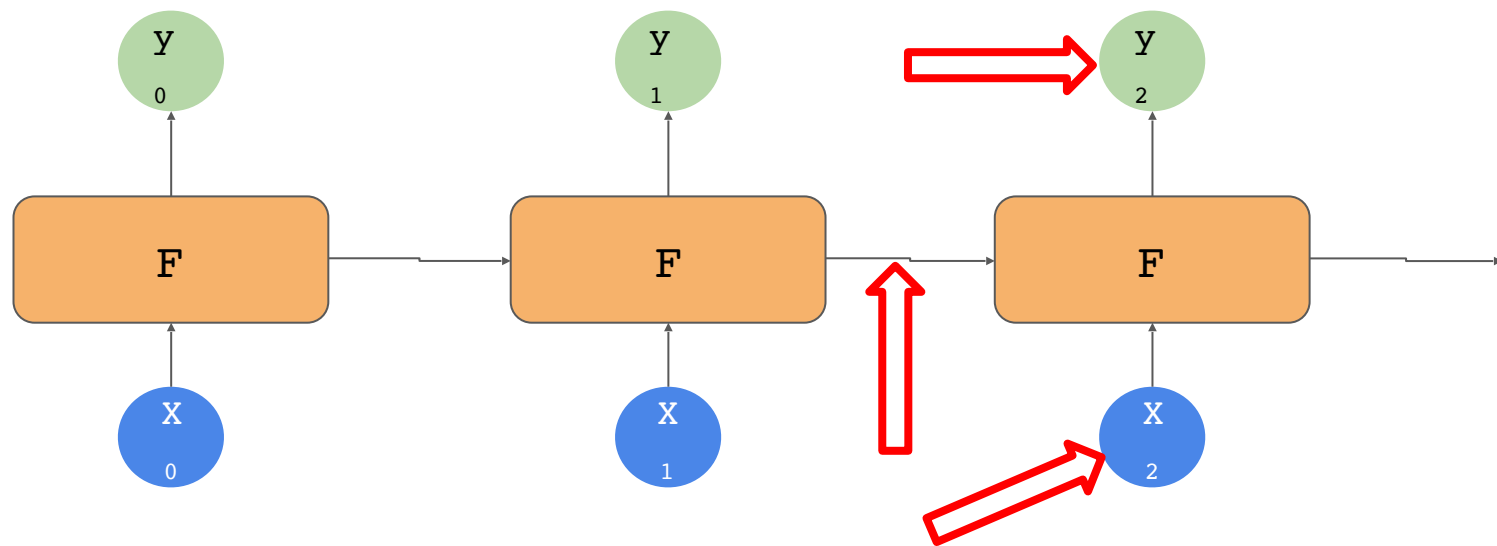





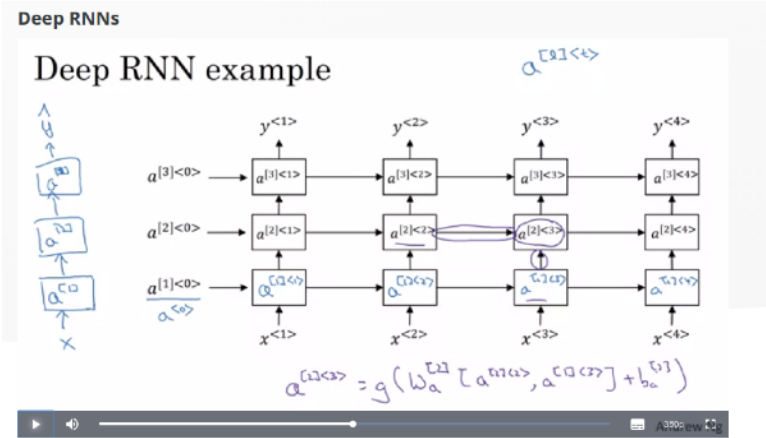












From the course by deeplearning.ai

Sequence Models

★★★★☆ 13393 ratings

Try the Course for Free

This Course Video Transcript

deeplearning.ai

Sequence Models

★★★★☆ 13393 ratings

Course 5 of 5 in the Specialization [Deep Learning](#)

This course will teach you how to build models for natural language, audio, and other sequence data. Thanks to deep learning, sequence algorithms are working far better than just two years ago, and this is enabling numerous exciting applications in speech recognition, music synthesis, chatbots, machine translation, natural language understanding, and many others. You will:

- Understand how to build and train Recurrent Neural Networks (RNNs), and commonly-

More

Today has a beautiful blue <...>

Today has a beautiful blue <...>

Today has a beautiful blue sky

Today has a beautiful blue <...>

Today has a beautiful blue sky




I lived in Ireland, so at school they made me learn how to speak <...>


I lived in Ireland, so at school they made me learn how to speak <...>

I lived in Ireland, so at school they made me learn how to speak Gaelic

I lived in Ireland, so at school they made me learn how to speak <...>

I lived in Ireland so at school they made me learn how to speak Gaelic






Long Short Term Memory (LSTM)

LSTM in pictures

$$\tilde{c}^{<t>} = \tanh(W_c[a^{<t-1>}, x^{<t>}] + b_c)$$
$$\Gamma_u = \sigma(W_u[a^{<t-1>}, x^{<t>}] + b_u)$$
$$\Gamma_f = \sigma(W_f[a^{<t-1>}, x^{<t>}] + b_f)$$
$$\Gamma_o = \sigma(W_o[a^{<t-1>}, x^{<t>}] + b_o)$$
$$c^{<t>} = \Gamma_u * \tilde{c}^{<t>} + \Gamma_f * c^{<t-1>}$$
$$a^{<t>} = \Gamma_o * \tanh c^{<t>}$$



From the course by deeplearning.ai

Sequence Models

★★★★☆ 13393 ratings

Try the Course for Free

This Course

Video Transcript



deeplearning.ai

Sequence Models

★★★★☆ 13393 ratings

Course 5 of 5 in the Specialization [Deep Learning](#)

This course will teach you how to build models for natural language, audio, and other sequence data. Thanks to deep learning, sequence algorithms are working far better than just two years ago, and this is enabling numerous exciting applications in speech recognition, music synthesis, chatbots, machine translation, natural language understanding, and many others. You will: - Understand how to build and train Recurrent Neural Networks (RNNs), and commonly-

More

```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(tokenizer.vocab_size, 64),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(64)),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

```
model = tf.keras.Sequential([  
    tf.keras.layers.Embedding(tokenizer.vocab_size, 64),  
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(64)),  
    tf.keras.layers.Dense(64, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])
```



```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(tokenizer.vocab_size, 64),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(64)),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

Layer (type)	Output Shape	Param #
=====		
embedding_2 (Embedding)	(None, None, 64)	523840
=====		
bidirectional_1 (Bidirectional)	(None, 128)	66048
=====		
dense_4 (Dense)	(None, 64)	8256
=====		
dense_5 (Dense)	(None, 1)	65
=====		
Total params: 598,209		
Trainable params: 598,209		
Non-trainable params: 0		

Layer (type)	Output Shape	Param #
=====		
embedding_2 (Embedding)	(None, None, 64)	523840
=====		
bidirectional_1 (Bidirection	(None, 128)	66048
=====		
dense_4 (Dense)	(None, 64)	8256
=====		
dense_5 (Dense)	(None, 1)	65
=====		
Total params: 598,209		
Trainable params: 598,209		
Non-trainable params: 0		


```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(tokenizer.vocab_size, 64),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(64, return_sequences=True)),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32)),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(tokenizer.vocab_size, 64),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(64, return_sequences=True)),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32)),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```


Layer (type)	Output Shape	Param #
=====		
embedding_3 (Embedding)	(None, None, 64)	523840
<hr/>		
bidirectional_2 (Bidirectional)	(None, None, 128)	66048
<hr/>		
bidirectional_3 (Bidirectional)	(None, 64)	41216
<hr/>		
dense_6 (Dense)	(None, 64)	4160
<hr/>		
dense_7 (Dense)	(None, 1)	65
=====		

Total params: 635,329  
Trainable params: 635,329  
Non-trainable params: 0

## 10 Epochs : Accuracy Measurement




1 Layer  
LSTM




2 Layer  
LSTM

## 10 Epochs : Loss Measurement




1 Layer  
LSTM




2 Layer  
LSTM



## 50 Epochs : Accuracy Measurement




1 Layer  
LSTM




2 Layer  
LSTM

## 50 Epochs : Loss Measurement



1 Layer  
LSTM




2 Layer  
LSTM


```
model = tf.keras.Sequential([  
    tf.keras.layers.Embedding(vocab_size, embedding_dim,  
                               input_length=max_length),  
    tf.keras.layers.GlobalAveragePooling1D(),  
    tf.keras.layers.Dense(24, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])
```

```
model = tf.keras.Sequential([  
    tf.keras.layers.Embedding(vocab_size, embedding_dim,  
                               input_length=max_length),  
    tf.keras.layers.GlobalAveragePooling1D(),  
    tf.keras.layers.Dense(24, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])
```


```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, embedding_dim,
                              input_length=max_length),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32)),
    tf.keras.layers.Dense(24, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```




Without  
LSTM



With LSTM



Without  
LSTM





With LSTM

```
model = tf.keras.Sequential([  
    tf.keras.layers.Embedding(vocab_size, embedding_dim,  
                               input_length=max_length),  
    tf.keras.layers.Conv1D(128, 5, activation='relu'),  
    tf.keras.layers.GlobalMaxPooling1D(),  
    tf.keras.layers.Dense(24, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])
```



```
model = tf.keras.Sequential([  
    tf.keras.layers.Embedding(vocab_size, embedding_dim,  
                               input_length=max_length),  
    tf.keras.layers.Conv1D(128, 5, activation='relu'),  
    tf.keras.layers.GlobalMaxPooling1D(),  
    tf.keras.layers.Dense(24, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])
```



```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, embedding_dim,
                              input_length=max_length),
    tf.keras.layers.Conv1D(128, 5, activation='relu'),
    tf.keras.layers.GlobalMaxPooling1D(),
    tf.keras.layers.Dense(24, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

max\_length = 120

tf.keras.layers.Conv1D(128, 5, activation='relu'),

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 120, 16)	16000
conv1d (Conv1D)	(None, 116, 128)	10368
global_max_pooling1d (Global Max Pooling1D)	(None, 128)	0
dense (Dense)	(None, 24)	3096
dense_1 (Dense)	(None, 1)	25

Total params: 29,489  
Trainable params: 29,489  
Non-trainable params: 0

```
max_length = 120
```

```
tf.keras.layers.Conv1D(128, 5, activation='relu'),
```

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 120, 16)	16000
conv1d (Conv1D)	(None, 116, 128)	0368
global_max_pooling1d (GlobalMaxPooling1D)	(None, 128)	0
dense (Dense)	(None, 24)	3096
dense_1 (Dense)	(None, 1)	25

Total params: 29,489  
Trainable params: 29,489  
Non-trainable params: 0

```
max_length = 120
```

```
tf.keras.layers.Conv1D(128, 5, activation='relu'),
```

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 120, 16)	16000
conv1d (Conv1D)	(None, 116, 128)	10368
global_max_pooling1d (GlobalMaxPooling1D)	(None, 128)	0
dense (Dense)	(None, 24)	3096
dense_1 (Dense)	(None, 1)	25

Total params: 29,489  
Trainable params: 29,489  
Non-trainable params: 0

```
imdb, info = tfds.load("imdb_reviews", with_info=True, as_supervised=True)
```

```
model = tf.keras.Sequential([  
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),  
    tf.keras.layers.Flatten(),  
    tf.keras.layers.Dense(6, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])
```



```
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
model.summary()
```

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 120, 16)	160000
flatten (Flatten)	(None, 1920)	0
dense (Dense)	(None, 6)	11526
dense_1 (Dense)	(None, 1)	7

Total params: 171,533  
Trainable params: 171,533  
Non-trainable params: 0





IMDB with Embedding-only : ~ 5s per epoch


```
imdb, info = tfds.load("imdb_reviews", with_info=True, as_supervised=True)
```

```
# Model Definition with LSTM
```


```
model = tf.keras.Sequential([  
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),  
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32)),  
    tf.keras.layers.Dense(6, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])  
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])  
model.summary()
```

Layer (type)	Output Shape	Param #
embedding_7 (Embedding)	(None, 120, 16)	160000
bidirectional_7 (Bidirection	(None, 64)	12544
dense_14 (Dense)	(None, 6)	390
dense_15 (Dense)	(None, 1)	7

Total params: 173,941  
Trainable params: 172,941  
Non-trainable params: 0



IMDB with LSTM ~43s per epoch



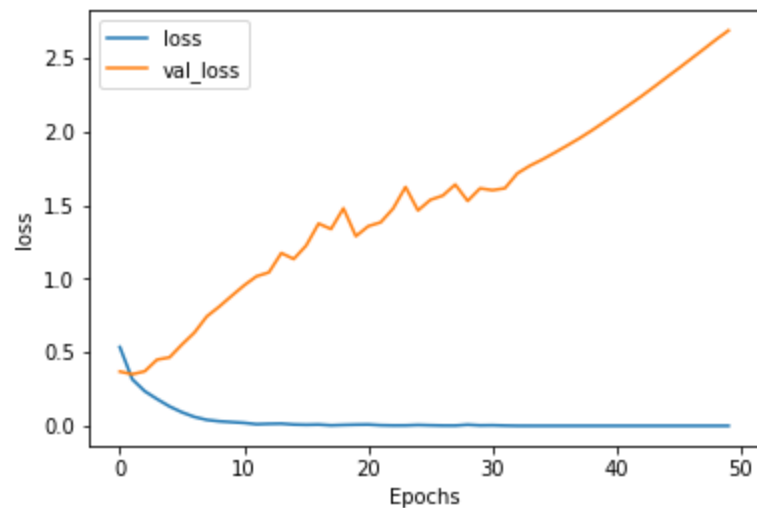

```
model = tf.keras.Sequential([  
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),  
    tf.keras.layers.Bidirectional(tf.keras.layers.GRU(32)),  
    tf.keras.layers.Dense(6, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])
```

```
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
model.summary()
```

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 120, 16)	160000
bidirectional_1 (Bidirectional)	(None, 64)	9600
dense_2 (Dense)	(None, 6)	390
dense_3 (Dense)	(None, 1)	7

Total params: 169,997  
Trainable params: 169,997  
Non-trainable params: 0





IMDB with GRU : ~ 20s per epoch

```
# Model Definition with Conv1D
```

```
model = tf.keras.Sequential([  
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),  
    tf.keras.layers.Conv1D(128, 5, activation='relu'),  
    tf.keras.layers.GlobalAveragePooling1D(),  
    tf.keras.layers.Dense(6, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
)  
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])  
model.summary()
```



Layer (type)	Output Shape	Param #
=====		
embedding (Embedding)	(None, 120, 16)	160000
<hr/>		
conv1d (Conv1D)	(None, 116, 128)	10368
<hr/>		
global_average_pooling1d (GlobalAveragePooling1D)	(None, 128)	0
<hr/>		
dense (Dense)	(None, 6)	774
<hr/>		
dense_1 (Dense)	(None, 1)	7
=====		
Total params: 171,149		
Trainable params: 171,149		
Non-trainable params: 0		
<hr/>		



IMDB with CNN : ~ 6s per epoch

# Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>

In the town of Athy one Jeremy Lanigan  
Battered away til he hadnt a pound.  
His father died and made him a man again  
Left him a farm and ten acres of ground.

He gave a grand party for friends and relations  
Who didnt forget him when come to the wall,  
And if youll but listen Ill make your eyes glisten  
Of the rows and the ructions of Lanigan's Ball.

Myself to be sure got free invitation,  
For all the nice girls and boys I might ask,  
And just in a minute both friends and relations  
Were dancing round merry as bees round a cask.

Judy ODaly, that nice little milliner,  
She tipped me a wink for to give her a call,  
And I soon arrived with Peggy McGilligan  
Just in time for Lanigans Ball.

```
tokenizer = Tokenizer()
```

```
data="In the town of Athy one Jeremy Lanigan \n Battered away ... ..."  
corpus = data.lower().split("\n")
```

```
tokenizer.fit_on_texts(corpus)
```

```
total_words = len(tokenizer.word_index) + 1
```

```
tokenizer = Tokenizer()
```

```
data="In the town of Athy one Jeremy Lanigan \n Battered away ... ..."
```

```
corpus = data.lower().split("\n")
```

```
tokenizer.fit_on_texts(corpus)
```

```
total_words = len(tokenizer.word_index) + 1
```

```
tokenizer = Tokenizer()
```

```
data="In the town of Athy one Jeremy Lanigan \n Battered away ... ..."
```

```
corpus = data.lower().split("\n")
```

```
tokenizer.fit_on_texts(corpus)
```

```
total_words = len(tokenizer.word_index) + 1
```

```
tokenizer = Tokenizer()
```

```
data="In the town of Athy one Jeremy Lanigan \n Battered away ... ..."
```

```
corpus = data.lower().split("\n")
```

```
tokenizer.fit_on_texts(corpus)
```

```
total_words = len(tokenizer.word_index) + 1
```



```
tokenizer = Tokenizer()
```

```
data="In the town of Athy one Jeremy Lanigan \n Battered away ... ..."
```

```
corpus = data.lower().split("\n")
```

```
tokenizer.fit_on_texts(corpus)
```

```
total_words = len(tokenizer.word_index) + 1
```

```
input_sequences = []  
for line in corpus:  
    token_list = tokenizer.texts_to_sequences([line])[0]  
    for i in range(1, len(token_list)):  
        n_gram_sequence = token_list[:i+1]  
        input_sequences.append(n_gram_sequence)
```

```
input_sequences = []
```

```
for line in corpus:
```

```
    token_list = tokenizer.texts_to_sequences([line])[0]
```

```
    for i in range(1, len(token_list)):
```

```
        n_gram_sequence = token_list[i+1]
```

```
        input_sequences.append(n_gram_sequence)
```

```
input_sequences = []  
for line in corpus:  
    token_list = tokenizer.texts_to_sequences([line])[0]  
    for i in range(1, len(token_list)):  
        n_gram_sequence = token_list[i+1]  
        input_sequences.append(n_gram_sequence)
```

In the town of Athy one Jeremy Lanigan



[4 2 66 8 67 68 69 70]

```
input_sequences = []  
for line in corpus:  
    token_list = tokenizer.texts_to_sequences([line])[0]  
    for i in range(1, len(token_list)):  
        n_gram_sequence = token_list[i+1:]  
        input_sequences.append(n_gram_sequence)
```

Line:

Input Sequences:

[4 2 66 8 67 68 69 70]

[4 2]

[4 2 66]

[4 2 66 8]

[4 2 66 8 67]

[4 2 66 8 67 68]

[4 2 66 8 67 68 69]

[4 2 66 8 67 68 69 70]

```
max_sequence_len = max([len(x) for x in input_sequences])
```



```
input_sequences =  
    np.array(pad_sequences(input_sequences, maxlen=max_sequence_len, padding='pre'))
```

Line:

Padded Input Sequences:

[4 2 66 8 67 68 69 70]

[0 0 0 0 0 0 0 0 0 4 2]

[0 0 0 0 0 0 0 0 0 4 2 66]

[0 0 0 0 0 0 0 0 4 2 66 8]

[0 0 0 0 0 0 0 4 2 66 8 67]

[0 0 0 0 0 0 4 2 66 8 67 68]

[0 0 0 0 0 4 2 66 8 67 68 69]

[0 0 0 0 4 2 66 8 67 68 69 70]

Padded Input Sequences:

[0 0 0 0 0 0 0 0 0 0 4 2]

[0 0 0 0 0 0 0 0 0 0 4 2 66]

[0 0 0 0 0 0 0 0 0 4 2 66 8]

[0 0 0 0 0 0 0 4 2 66 8 67]

[0 0 0 0 0 0 4 2 66 8 67 68]

[0 0 0 0 0 4 2 66 8 67 68 69]

[0 0 0 0 4 2 66 8 67 68 69 70]

Padded Input Sequences:

Input (X)

Label (Y)

[0 0 0 0 0 0 0 0 0 4 2]

[0 0 0 0 0 0 0 0 0 4 2 66]

[0 0 0 0 0 0 0 0 4 2 66 8]

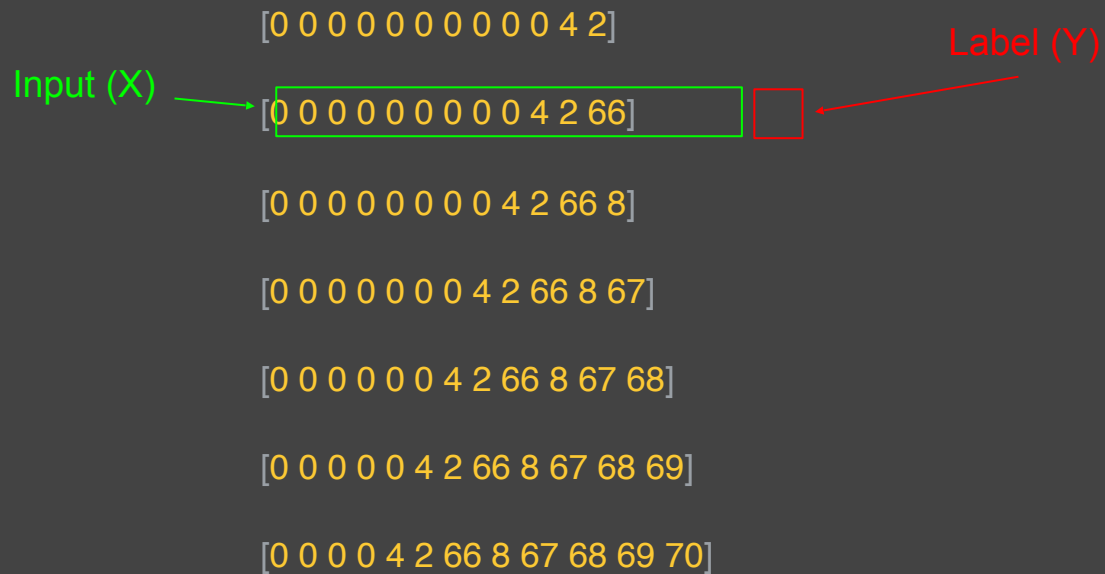
[0 0 0 0 0 0 0 4 2 66 8 67]

[0 0 0 0 0 0 4 2 66 8 67 68]

[0 0 0 0 0 4 2 66 8 67 68 69]

[0 0 0 0 4 2 66 8 67 68 69 70]

Padded Input Sequences:



Padded Input Sequences:

[0 0 0 0 0 0 0 0 0 4 2]

[0 0 0 0 0 0 0 0 0 4 2 66]

Input (X)

[0 0 0 0 0 0 0 0 4 2 66 8]

Label (Y)



[0 0 0 0 0 0 0 4 2 66 8 67]

[0 0 0 0 0 0 4 2 66 8 67 68]

[0 0 0 0 0 4 2 66 8 67 68 69]

[0 0 0 0 4 2 66 8 67 68 69 70]

```
xs = input_sequences[:, :-1]  
labels = input_sequences[:, -1]
```

```
ys = tf.keras.utils.to_categorical(labels, num_classes=total_words)
```







```
model = Sequential()
model.add(Embedding(total_words, 64, input_length=max_sequence_len - 1))
model.add((LSTM(20)))
model.add(Dense(total_words, activation='softmax'))
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
model.fit(xs, ys, epochs=500, verbose=1)
```

```
model = Sequential()
model.add(Embedding(total_words, 64, input_length=max_sequence_len - 1))
model.add((LSTM(20)))
model.add(Dense(total_words, activation='softmax'))
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
model.fit(xs, ys, epochs=500, verbose=1)
```

```
model = Sequential()
model.add(Embedding(total_words, 64, input_length=max_sequence_len - 1))
model.add(LSTM(20))
model.add(Dense(total_words, activation='softmax'))
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
model.fit(xs, ys, epochs=500, verbose=1)
```


```
model = Sequential()
model.add(Embedding(total_words, 64, input_length=max_sequence_len - 1))
model.add(LSTM(20))
model.add(Dense(total_words, activation='softmax'))
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
model.fit(xs, ys, epochs=500, verbose=1)
```

```
model = Sequential()
model.add(Embedding(total_words, 64, input_length=max_sequence_len - 1))
model.add((LSTM(20)))
model.add(Dense(total_words, activation='softmax'))
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
model.fit(xs, ys, epochs=500, verbose=1)
```

```
model = Sequential()
model.add(Embedding(total_words, 64, input_length=max_sequence_len - 1))
model.add((LSTM(20)))
model.add(Dense(total_words, activation='softmax'))
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
model.fit(xs, ys, epochs=500, verbose=1)
```



```
model = Sequential()
model.add(Embedding(total_words, 64, input_length=max_sequence_len - 1))
model.add((LSTM(20)))
model.add(Dense(total_words, activation='softmax'))
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
model.fit(xs, ys, epochs=500, verbose=1)
```



Laurence went to dublin round the plenty as red wall me for wall wall

Laurence went to dublin odaly of the nice of lanigans ball ball ball hall

Laurence went to dublin he hadnt a minute both relations hall new relations youd

Laurence went to dublin round the plenty as red wall me for wall wall

Laurence went to dublin odaly of the nice of lanigans ball ball ball hall

Laurence went to dublin he hadnt a minute both relations hall new relations youd

Laurence went to dublin round the plenty as red wall me for wall wall

Laurence went to dublin odaly of the nice of lanigans ball ball ball hall



Laurence went to dublin he hadnt a minute both relations hall new relations youd

Laurence went to dublin round the plenty as red wall me for wall wall

Laurence went to dublin odaly of the nice of lanigans ball ball ball hall

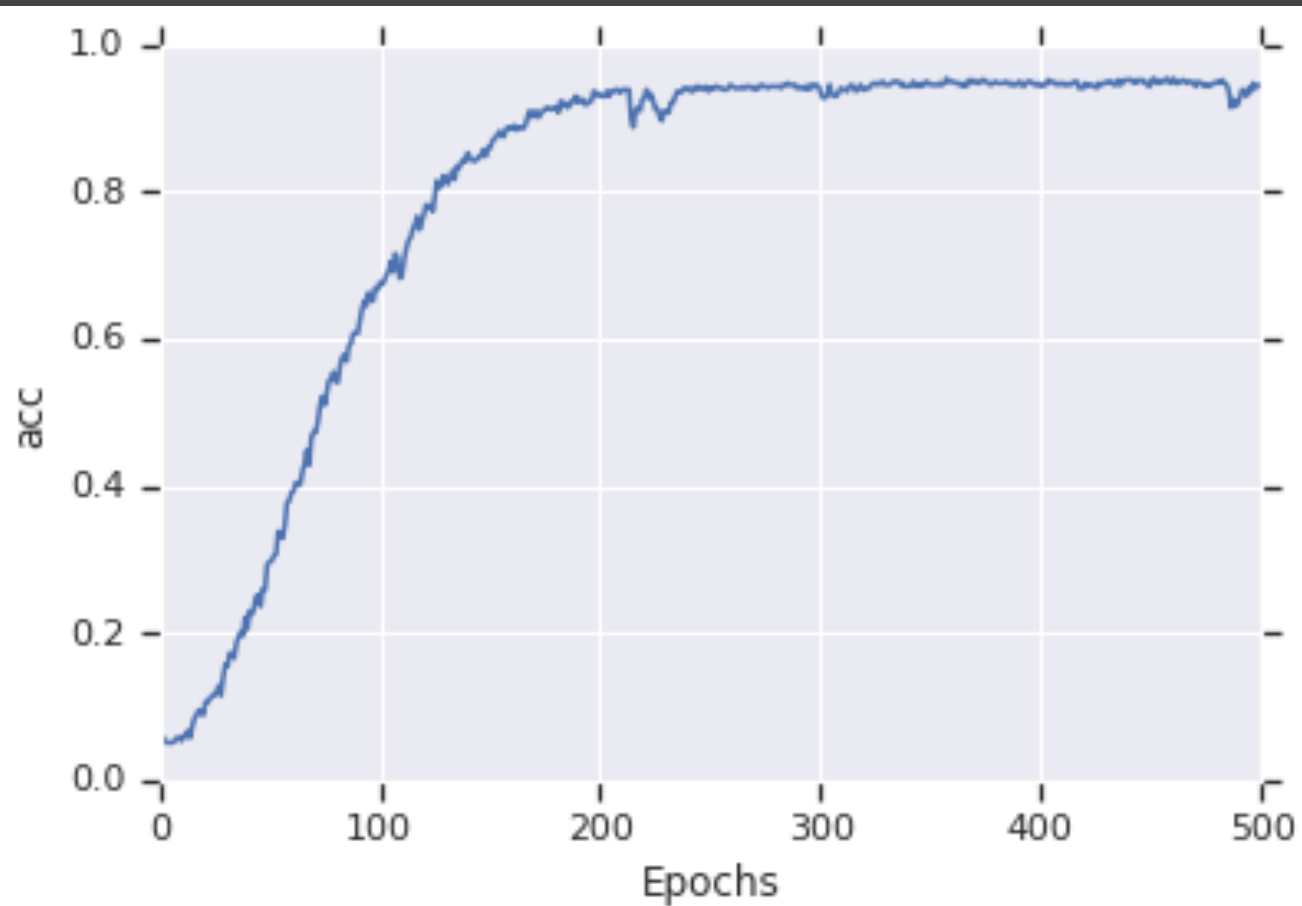
Laurence went to dublin he hadnt a minute both relations hall new relations you d



```
model = Sequential()  
model.add(Embedding(total_words, 64, input_length=max_sequence_len - 1))  
model.add(Bidirectional(LSTM(20)))  
model.add(Dense(total_words, activation='softmax'))  
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])  
model.fit(xs, ys, epochs=500, verbose=1)
```

```
model = Sequential()
model.add(Embedding(total_words, 64, input_length=max_sequence_len - 1))
model.add(Bidirectional(LSTM(20)))
model.add(Dense(total_words, activation='softmax'))
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
model.fit(xs, ys, epochs=500, verbose=1)
```






Laurence went to dublin think and wine for lanigans ball entangled in nonsense me  
Laurence went to dublin his pipes bellows chanter and all all entangled all kinds  
Laurence went to dublin how the room a whirligig ructions long at brooks fainted

Laurence went to dublin

```
token_list = tokenizer.texts_to_sequences([seed_text])[0]
```

Laurence went to dublin

[134, 13, 59]



```
token_list = pad_sequences([token_list], maxlen=max_sequence_len - 1, padding='pre')
```



```
predicted = model.predict(token_list)
predicted = np.argmax(probabilities, axis=-1)[0]
```



```
output_word = tokenizer.index_word[predicted]  
seed_text += " " + output_word
```

```
seed_text = "Laurence went to dublin"
```

```
next_words = 10
```

```
for _ in range(next_words):
```

```
    token_list = tokenizer.texts_to_sequences([seed_text])[0]
```

```
    token_list = pad_sequences([token_list], maxlen=max_sequence_len - 1, padding='pre')
```

```
    predicted = model.predict_classes(token_list, verbose=0)
```

```
    output_word = output_word = tokenizer.index_word[predicted]
```

```
    seed_text += " " + output_word
```

```
print(seed_text)
```

[illegible]

```
!wget --no-check-certificate \  
  https://storage.googleapis.com/laurencemoroney-blog.appspot.com/irish-lyrics-eof.txt \  
  -O /tmp/irish-lyrics-eof.txt
```

```
data = open('/tmp/irish-lyrics-eof.txt').read()
```

```
model = Sequential()
model.add(Embedding(total_words, 100, input_length=max_sequence_len-1))
model.add(Bidirectional(LSTM(150)))
model.add(Dense(total_words, activation='softmax'))
adam = Adam(lr=0.01)
model.compile(loss='categorical_crossentropy', optimizer=adam, metrics=['accuracy'])
history = model.fit(xs, ys, epochs=100, verbose=1)
```

```
model = Sequential()
model.add(Embedding(total_words, 100, input_length=max_sequence_len-1))
model.add(Bidirectional(LSTM(150)))
model.add(Dense(total_words, activation='softmax'))
adam = Adam(lr=0.01)
model.compile(loss='categorical_crossentropy', optimizer=adam, metrics=['accuracy'])
history = model.fit(xs, ys, epochs=100, verbose=1)
```

```
model = Sequential()
model.add(Embedding(total_words, 100, input_length=max_sequence_len-1))
model.add(Bidirectional(LSTM(150)))
model.add(Dense(total_words, activation='softmax'))
adam = Adam(lr=0.01)
model.compile(loss='categorical_crossentropy', optimizer=adam, metrics=['accuracy'])
history = model.fit(xs, ys, epochs=100, verbose=1)
```



```
model = Sequential()
model.add(Embedding(total_words, 100, input_length=max_sequence_len-1))
model.add(Bidirectional(LSTM(150)))
model.add(Dense(total_words, activation='softmax'))
adam = Adam(lr=0.01)
model.compile(loss='categorical_crossentropy', optimizer=adam, metrics=['accuracy'])
history = model.fit(xs, ys, epochs=100, verbose=1)
```

```
model = Sequential()
model.add(Embedding(total_words, 100, input_length=max_sequence_len-1))
model.add(Bidirectional(LSTM(150)))
model.add(Dense(total_words, activation='softmax'))
adam = Adam(lr=0.01)
model.compile(loss='categorical_crossentropy', optimizer=adam, metrics=['accuracy'])
history = model.fit(xs, ys, epochs=100, verbose=1)
```

Help Me Obi-Wan Kenobi, you're my only hope  
my dear  
and hope as i did fly with its flavours  
along with all its joys  
but sure i will build  
love you still  
gold it did join  
do mans run away cross our country  
are wedding i was down to  
off holyhead wished meself  
down among the pigs  
played some hearty rigs  
me embarrass  
find me brother  
me chamber she gave me  
who storied be irishmen  
to greet you  
lovely molly  
gone away from me home  
home to leave the old tin cans  
the foemans chain one was shining  
sky above i think i love

