

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>



deeplearning.ai

Generative Models


Outline

- What are generative models
- Types of generative models



Generative Models vs. Discriminative Models

Discriminative models



Features Class

$$X \rightarrow Y$$

$$P(Y|X)$$

Generative models



Noise Class Features

$$\xi, Y \rightarrow X$$

$$P(X|Y)$$

Generative Models vs. Discriminative Models



Generative models




Noise Class Features

$$\xi, Y \rightarrow X$$

$$P(X|Y)$$

Generative Models

Variational Autoencoders

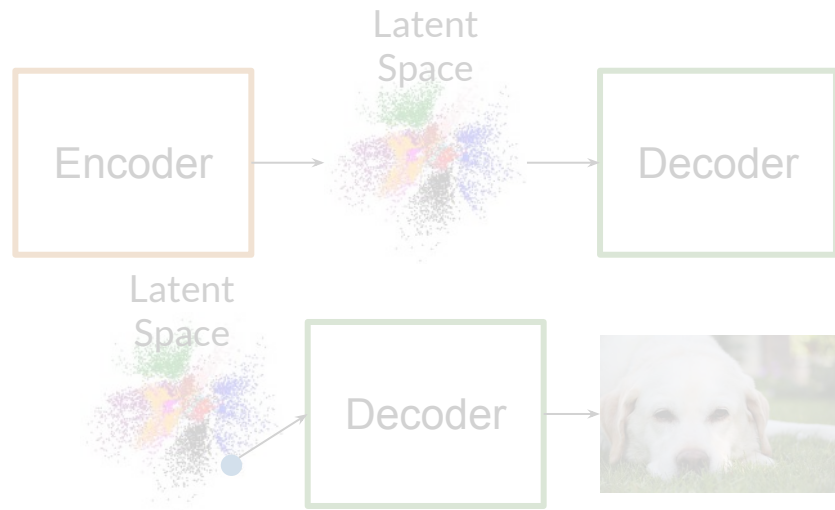


Generative Adversarial Networks


Available from: <https://arxiv.org/abs/1804.00891>

Generative Models

Variational Autoencoders




Generative Adversarial Networks



Available from: <https://arxiv.org/abs/1804.00891>

Summary

- Generative models learn to produce examples
- Discriminative models distinguish between classes
- Up next, GANs!






deeplearning.ai

Real Life GANs

Outline

- Cool applications of GANs
- Major companies using them



GANs Over Time



Ian Goodfellow
@goodfellow_ian




4.5 years of GAN progress on face generation.

arxiv.org/abs/1406.2661 arxiv.org/abs/1511.06434

arxiv.org/abs/1606.07536 arxiv.org/abs/1710.10196

arxiv.org/abs/1812.04948



GANs Over Time



Face Generation
StyleGAN2

These people do
not exist!

Karras, Tero, et al. "Analyzing and improving the image quality of stylegan." *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2020.

GANs Over Time



StyleGAN2



Mimics the
distribution of the
training data

Karras, Tero, et al. "Analyzing and improving the image quality of stylegan." *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2020.

<https://9gag.com/gag/aWYZKWx>

GANs for Image Translation


From one domain to another

CycleGAN



Park, Taesung, et al. "Semantic image synthesis with spatially-adaptive normalization." *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*. 2019.

GANs for Image Translation



GauGAN

Doodles



Pictures

Park, Taesung, et al. "Semantic image synthesis with spatially-adaptive normalization." *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*. 2019.

GANs are Magic!



Zakharov, Egor, et al. "Few-shot adversarial learning of realistic neural talking head models." *Proceedings of the IEEE International Conference on Computer Vision*. 2019.


GANs for 3D Objects

Chair




3D-GAN


Car



Sofa



Table



Generative
Design

Wu, Jiajun, et al. "Learning a probabilistic latent space of object shapes via 3d generative-adversarial modeling." *Advances in neural information processing systems*. 2016.

Companies Using GANs



Next-gen
Photoshop



Text
Generation



Data
Augmentation





Image Filters



Super-resolution

Summary

- GANs' performance is rapidly improving
- Huge opportunity to work in this space!
- Major companies are using them





deeplearning.ai

Intuition Behind GANs

Outline

- The goal of the generator and the discriminator
- The competition between them




Generative Adversarial Network

Generator learns to make *fakes*
that look **real**



Discriminator learns to distinguish
real from *fake*



Generative Adversarial Network

Discriminator learns to distinguish
real from *fake*



Generative Adversarial Network

Generator learns to make *fakes*
that look **real**

Discriminator learns to distinguish
real from *fake*



Generative Adversarial Network

Generator learns to make *fakes*
that look *real*



Doesn't know how
it should look



Generative Adversarial Network

Discriminator learns to distinguish
real from *fake*



The Game Is On!



5% Real




40% Real



80% Real




The Game is On!



The Game Is On!



30% Real



60% Real




95% Real

The Game Is On!




30% Real



60% Real




Wrong!



95% Real

Won't fool me
again

The Game Is On!



Summary

- The **generator's** goal is to fool the discriminator
- The **discriminator's** goal is to distinguish between real and fake
- They learn from the competition with each other
- At the end, *fakes* look **real**






deeplearning.ai

Discriminator

Outline

- Review of classifiers
- The role of classifiers in terms of probability
- Discriminator



Classifiers

Distinguish between different classes



Turtle

Bird


Cat

Dog


Fish

Classifiers


Distinguish between different classes



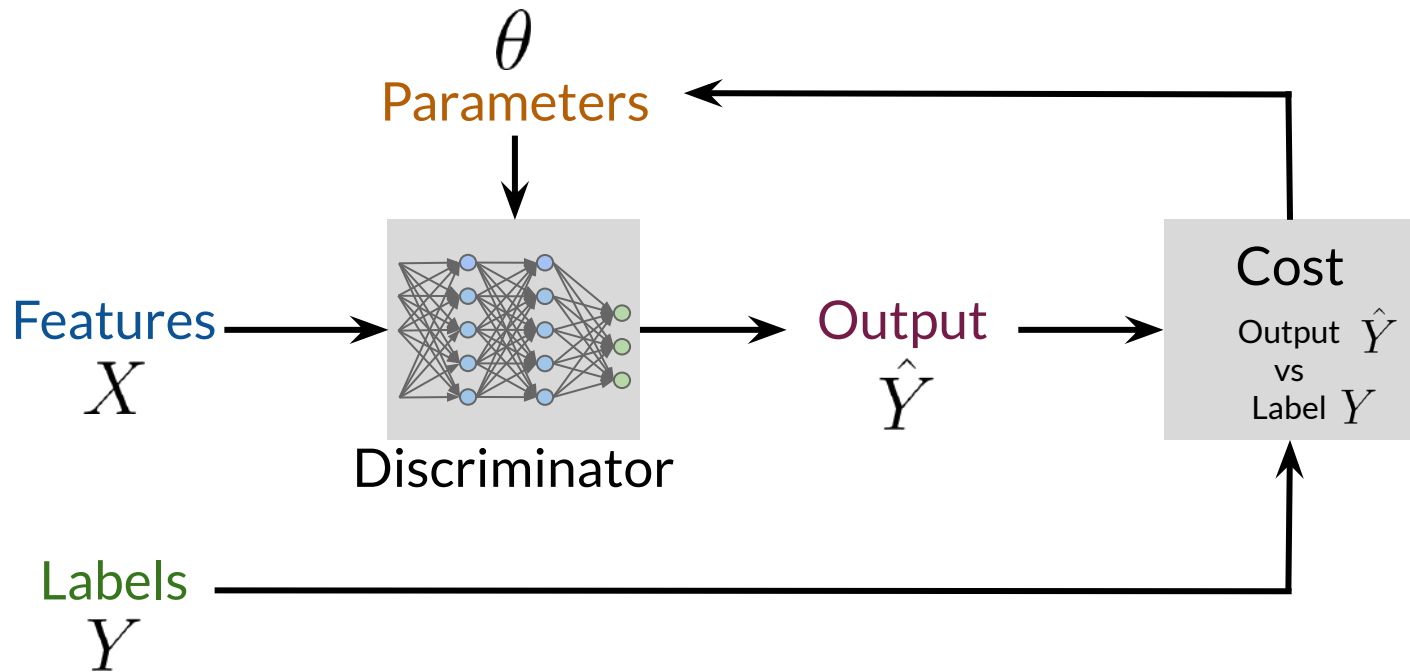
Neural Networks



Neural Networks



Classifiers (training)




Classifiers



$$P\left(\begin{array}{c} \text{Turtle} \\ \text{Bird} \\ \text{Cat} \\ \text{Dog} \\ \text{Fish} \end{array} \mid \text{Image} \right)$$





Classifiers



Discriminator



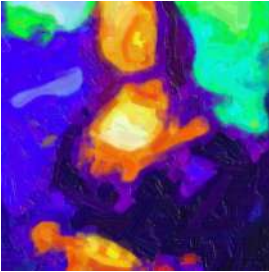
Discriminator



Discriminator

$$P\left(\begin{array}{c} \text{Fake} \\ \text{Class} \end{array} \mid \begin{array}{c} X \\ \text{Features} \end{array} \right)$$

Discriminator

$$P\left(\begin{array}{c} \text{Fake} \\ \text{Class} \end{array} \mid \begin{array}{c} \text{Features} \end{array}\right) = 0.85 \rightarrow \boxed{\text{Fake}}$$


Summary

- The **discriminator** is a classifier
- It learns the probability of class Y (**real** or **fake**) given features X
- The probabilities are the feedback for the **generator**





deeplearning.ai

Generator

Outline

- What the generator does
- How it improves its performance
- Generator in terms of probability



Generator

Turtle

Generates examples of the class

Bird


Cat

Dog


Fish



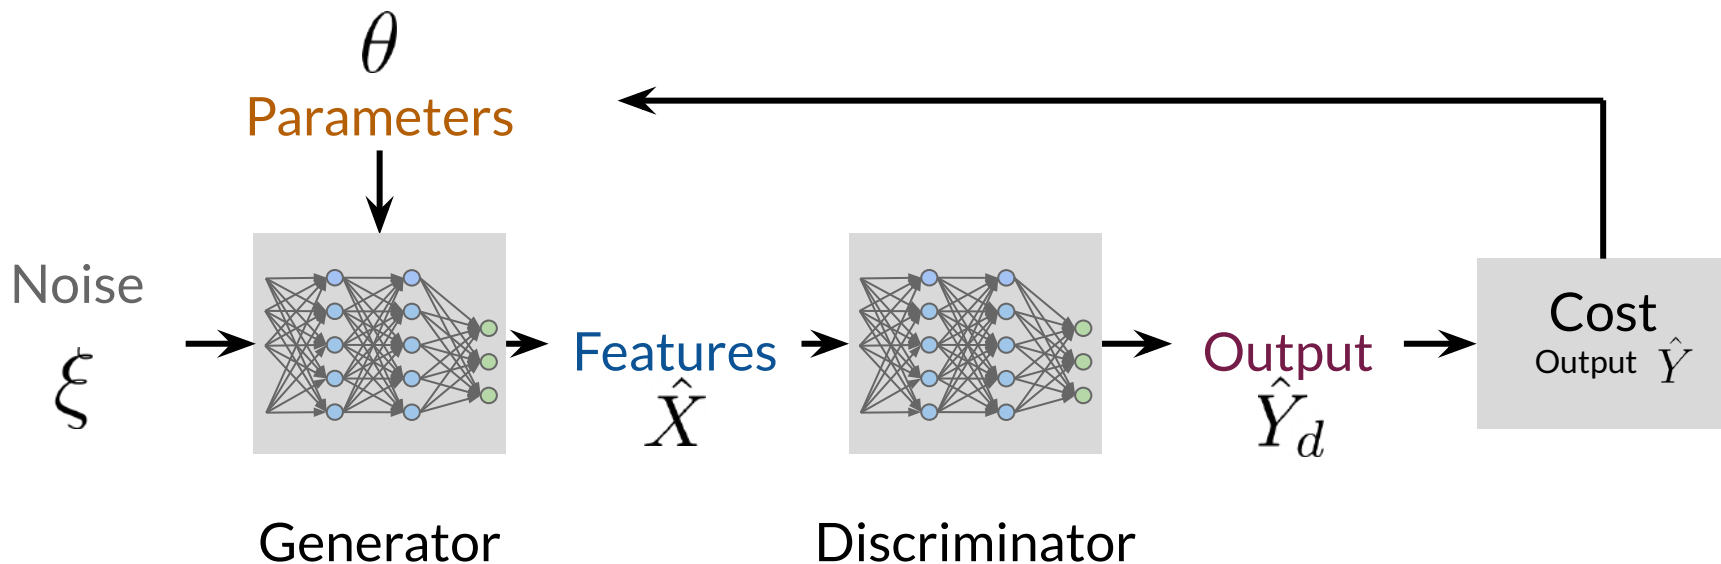
Neural Networks




Neural Networks



Generator: Learning



Sampling



Generator

$$P(\text{Image} \mid \text{Class})$$

Turtle


Bird

Cat

Dog

Fish

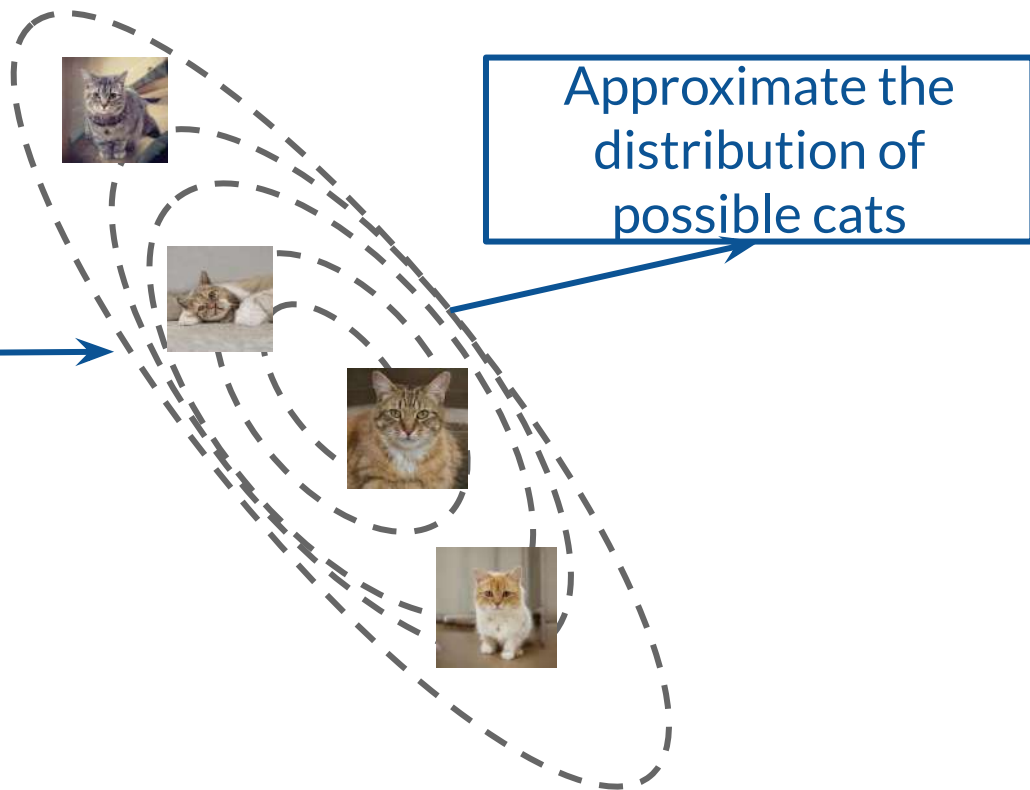
Generator



Generator

$P(X)$

Features



Images available from: <http://thesecatsdonotexist.com/>

Summary

- The **generator** produces fake data
- It learns the probability of features X
- The **generator** takes as input noise (random features)





deeplearning.ai

BCE Cost Function

Outline

- Binary Cross Entropy (BCE) Loss equation by parts
- How it looks graphically



BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

Prediction

Label

Features

Parameters

Average loss of the whole batch

The diagram illustrates the components of the Binary Cross-Entropy (BCE) cost function. The formula is shown as $J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$. Annotations include: a circle around the summation term with an arrow pointing to a box labeled 'Average loss of the whole batch'; an arrow from $h(x^{(i)}, \theta)$ pointing to 'Prediction'; an arrow from $y^{(i)}$ pointing to 'Label'; an arrow from $x^{(i)}$ pointing to 'Features'; and an arrow from θ pointing to 'Parameters'.

BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

$y^{(i)}$	$h(x^{(i)}, \theta)$	$y^{(i)} \log h(x^{(i)}, \theta)$
0	any	0
1	0.99	~ 0
1	~ 0	-inf

Relevant when
the label is 1

BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

$y^{(i)}$	$h(x^{(i)}, \theta)$	$(1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))$
1	any	0
0	0.01	~ 0
0	~ 1	-inf

Relevant when
the label is 0


BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

Ensures that the cost is always greater or equal to 0



BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$



BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$



Summary

- The BCE cost function has two parts (one relevant for each class)
- Close to zero when the label and the prediction are similar
- Approaches infinity when the label and the prediction are different






deeplearning.ai


Putting It All Together

Outline

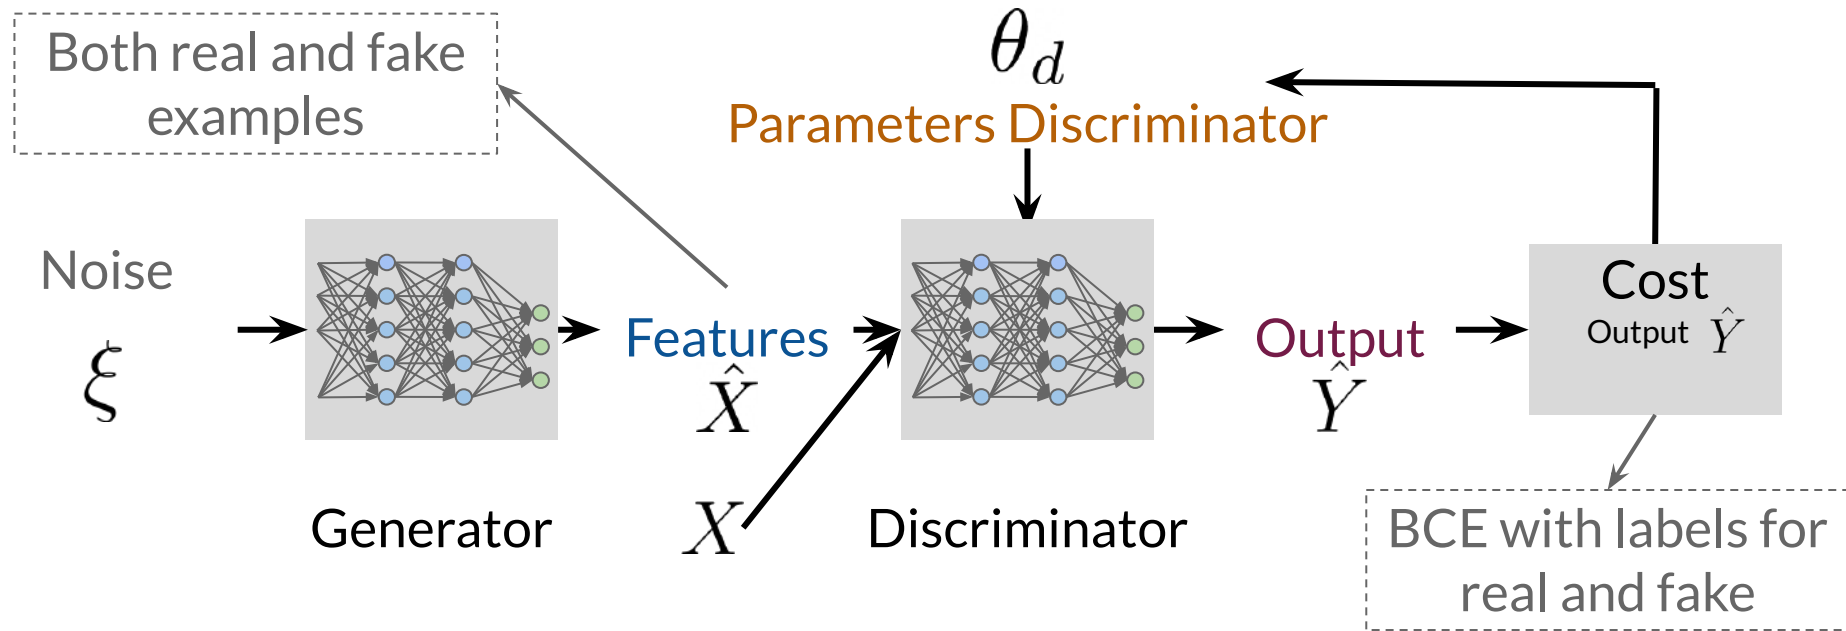
- How the whole architecture looks
- How to train GANs



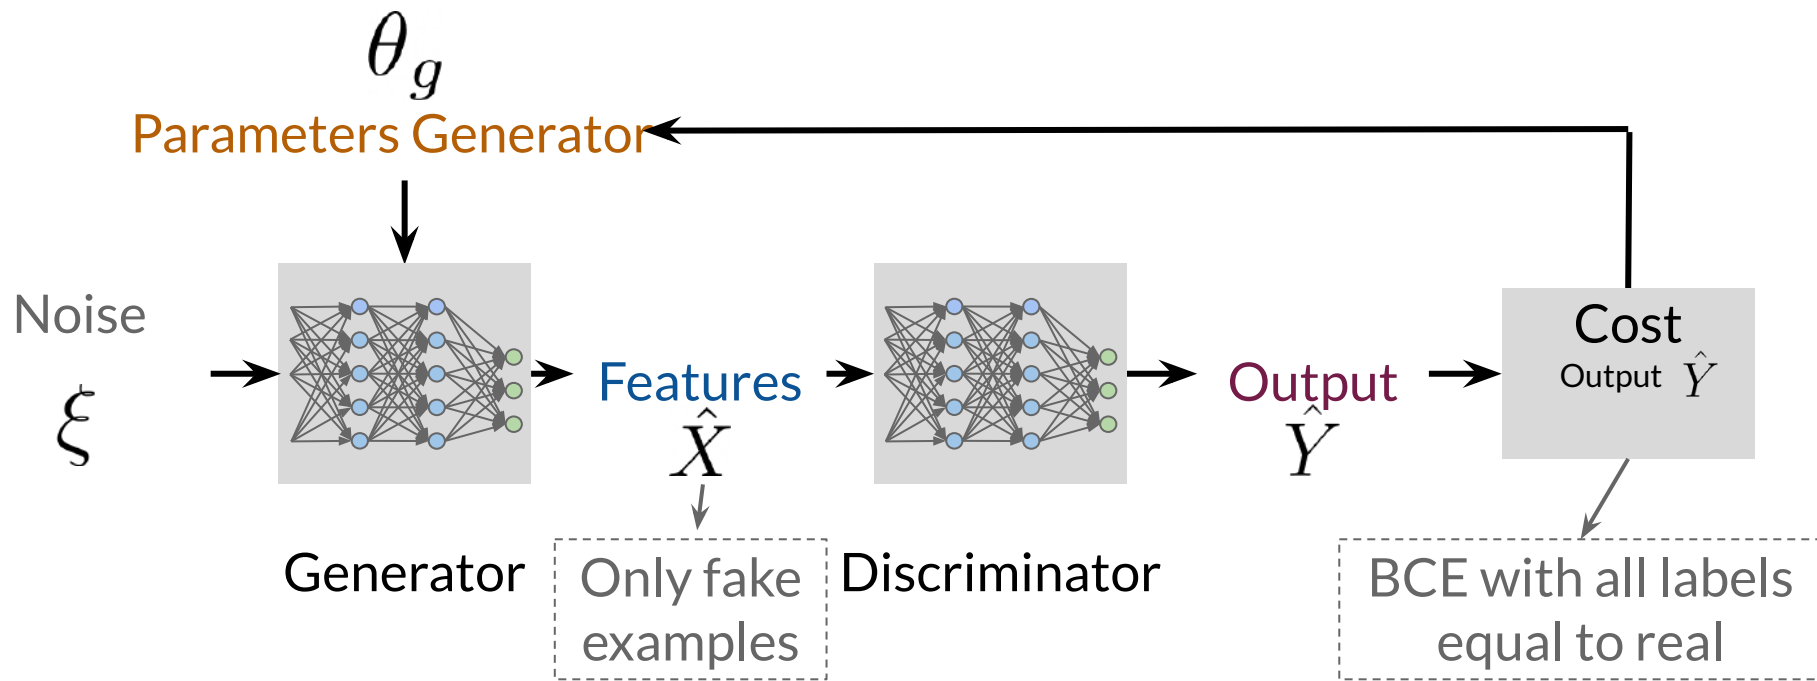
GANs Model



Training GANs: Discriminator



Training GANs: Generator




Training GANs



Summary

- GANs train in an alternating fashion
- The two models should always be at a similar “skill” level





deeplearning.ai

Intro to PyTorch (Optional)

Outline

- Comparison with TensorFlow
- Defining Models
- Training



PyTorch vs TensorFlow

PyTorch

Imperative, computations on the go

```
A, B = 1, 2
C = A + B
print(C)
```

3

Dynamic Computational Graphs

TensorFlow

Symbolic, first define and then compile

```
C = A + B
f = compile(C)
print(f(A = 1, B = 2))
```

3

Static Computational Graphs

Tensorflow > 2.0 moves toward PyTorch by including **Eager Execution**

PyTorch vs TensorFlow

PyTorch

TensorFlow

Currently very similar frameworks!

Tensorflow > 2.0 moves toward PyTorch by
including Eager Execution

Defining Models in PyTorch

```
import torch
from torch import nn
```

→ Custom layers for DL

```
class LogisticRegression(nn.Module):
    def __init__(self, in):
        super().__init__()
        self.log_reg = nn.Sequential(
            nn.Linear(in, 1),
            nn.Sigmoid()
        )
    def forward(self, x):
        return self.log_reg(x)
```

Define the model as a class

Initialization method with parameters

Definition of the architecture

Forward computation of the model
with inputs x

Training Models In PyTorch

```
model = LogisticRegression(16)
```

Initialization of the model

```
criterion = nn.BCELoss()
```

Cost function

```
optimizer = torch.optim.SGD(model.parameters(), lr=0.01)
```

Optimizer

```
for t in range(n_epochs):
```

Training loop for number of epochs

```
    y_pred = model(x)
```

```
    loss = criterion(y_pred, y)
```

Forward propagation

```
    optimizer.zero_grad()
```

```
    loss.backward()
```

```
    optimizer.step()
```

Optimization step

Summary

- PyTorch makes computations on the run
- Dynamic computational graphs in Pytorch
- Just another framework, and similar to Tensorflow!



Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>




deeplearning.ai



Activations (Basic Properties)

Outline

- What are activations
- Reasoning behind non-linear differential activations



Activations



$$z_i^{[l]} = \sum_{i=0} W_i^{[l]} a_i^{[l-1]}$$

$$a_i^{[l]} = \boxed{g^{[l]}}(z_i^{[l]})$$

Differentiable
non-linear
function


Activations

$$a_i^{[l]} = \boxed{g^{[l]}}(z_i^{[l]})$$

Differentiable
non-linear
function

1. Differentiable for backpropagation

2. Non-linear to compute complex features, **if not**:



\equiv

$$WX + b$$

Linear
regression

Summary

- Activation functions are non-linear and differentiable
- Differentiable for backpropagation
- Non-linear to approximate complex functions





deeplearning.ai

Common Activation Functions

Outline


- Common activations and their structure
 - ReLU
 - Leaky ReLU
 - Sigmoid
 - Tanh



Activations: ReLU


ReLU = Rectified Linear
Unit

$$g^{[l]}(z^{[l]}) = \max(\underline{0}, \underline{z^{[l]}})$$



Dying ReLU
problem

Activations: Leaky ReLU




$$g^{[l]}(z^{[l]}) = \max(\underline{\alpha z^{[l]}}, \underline{z^{[l]}})$$

Solves the dying
ReLU problem

Activations: Sigmoid

$$g^{[l]}(z^{[l]}) = \frac{1}{1 + e^{-z^{[l]}}}$$




Values between 0
and 1

Vanishing gradient
and saturation
problems

Activations: Tanh

$$g^{[l]}(z^{[l]}) = \tanh(z^{[l]})$$




Values between -1
and 1

Keeps the sign of
the input

Same issues as
Sigmoid

Summary

- ReLU activations suffer from dying ReLU
- Leaky ReLU solve the dying ReLU problem
- Sigmoid and Tanh have vanishing gradient and saturation problems





deeplearning.ai

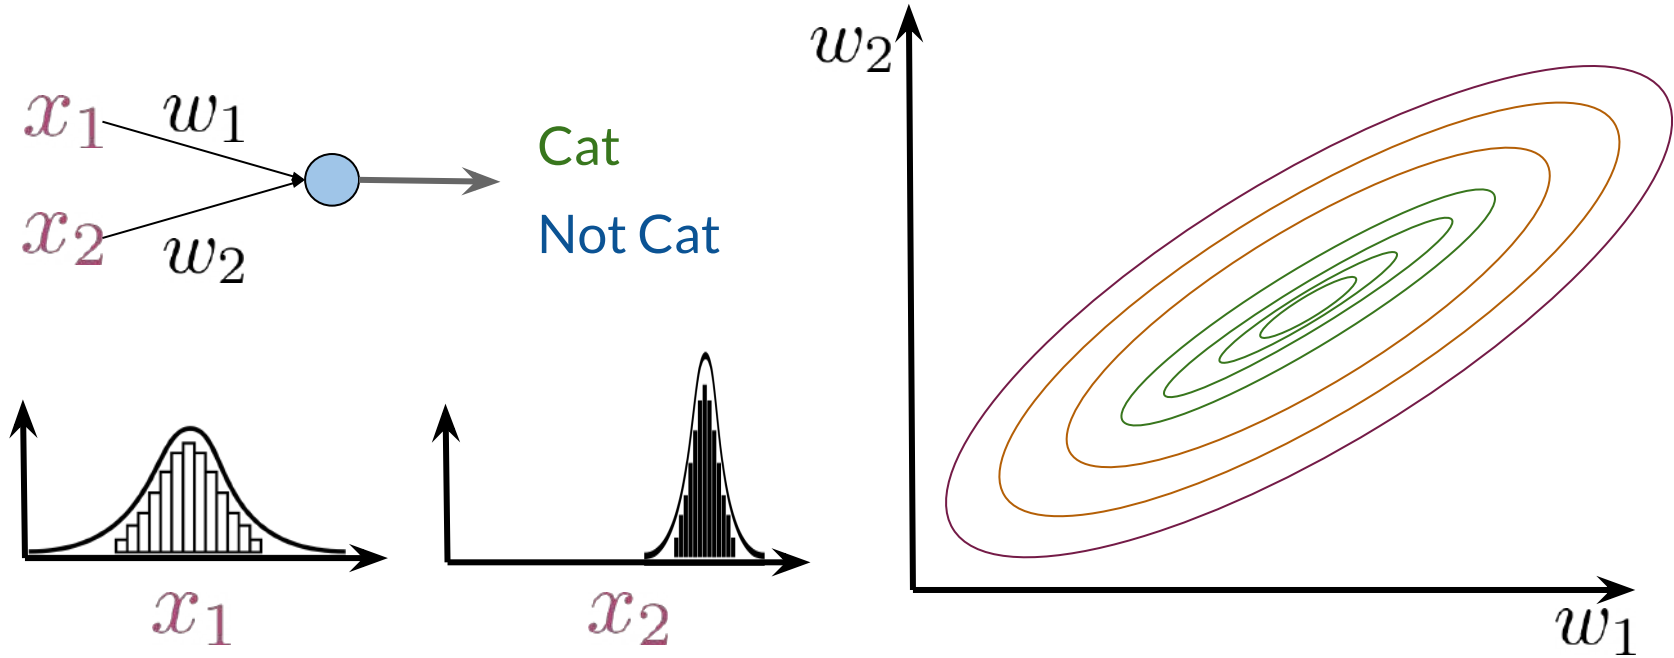
Batch Normalization (Explained)

Outline


- How normalization helps models
- Internal covariate shift
- Batch normalization




Different Distributions




Covariate Shift



Normalization and Its Effects



Normalization and Its Effects




x_1' x_2'

Around mean at 0
and std. at 1

Training data uses
batch stats

Test data uses
training stats

Normalization and Its Effects




x_1' x_2'

Around mean at 0
and std. at 1

Reduction of
covariate shift

Internal Covariate Shift




Changes in
weights




Changes in
activation
distribution

Batch Normalization




Normalizes the
input for each
neuron



Summary

- Batch normalization smooths the cost function
- Batch normalization reduces the internal covariate shift
- Batch normalization speeds up learning!






deeplearning.ai



Batch Normalization (Procedure)

Outline

- Batch norm for training
- Batch norm for testing



Batch Normalization: Training





$$z_i^{[l]} = \sum_{i=0} W_i^{[l]} a_i^{[l-1]} \rightarrow \text{For every example in the batch}$$

$$\hat{z}_i^{[l]} = \frac{z_i^{[l]} - \mu_{z_i^{[l]}}}{\sqrt{\sigma_{z_i^{[l]}}^2 + \epsilon}}$$

Batch mean of $z_i^{[l]}$
Batch std of $z_i^{[l]}$

Batch Normalization: Training



$$\hat{z}_i^{[l]} = \frac{z_i^{[l]} - \mu_{z_i^{[l]}}}{\sqrt{\sigma_{z_i^{[l]}}^2 + \epsilon}}$$


$$y_i^{[l]} = \boxed{\gamma} \hat{z}_i^{[l]} + \boxed{\beta}$$

Shift factor

Scale Factor

Learnable parameters to get the optimal dist.

Batch Normalization: Test



$$\hat{z}_i^{[l]} = \frac{z_i^{[l]} - \mathbf{E}(z_i^{[l]})}{\sqrt{\text{Var}(z_i^{[l]}) + \epsilon}}$$

Running **mean** and running **std** from training

Frameworks like
Tensorflow and Pytorch
keep track of them

Summary

- Batch norm introduces learnable shift and scale factors
- During test, the running statistics from training are used
- Frameworks take care of the whole process






deeplearning.ai


Review of Convolutions

Outline


- What convolutions are
- How they work




What is a convolution?




What is a convolution?



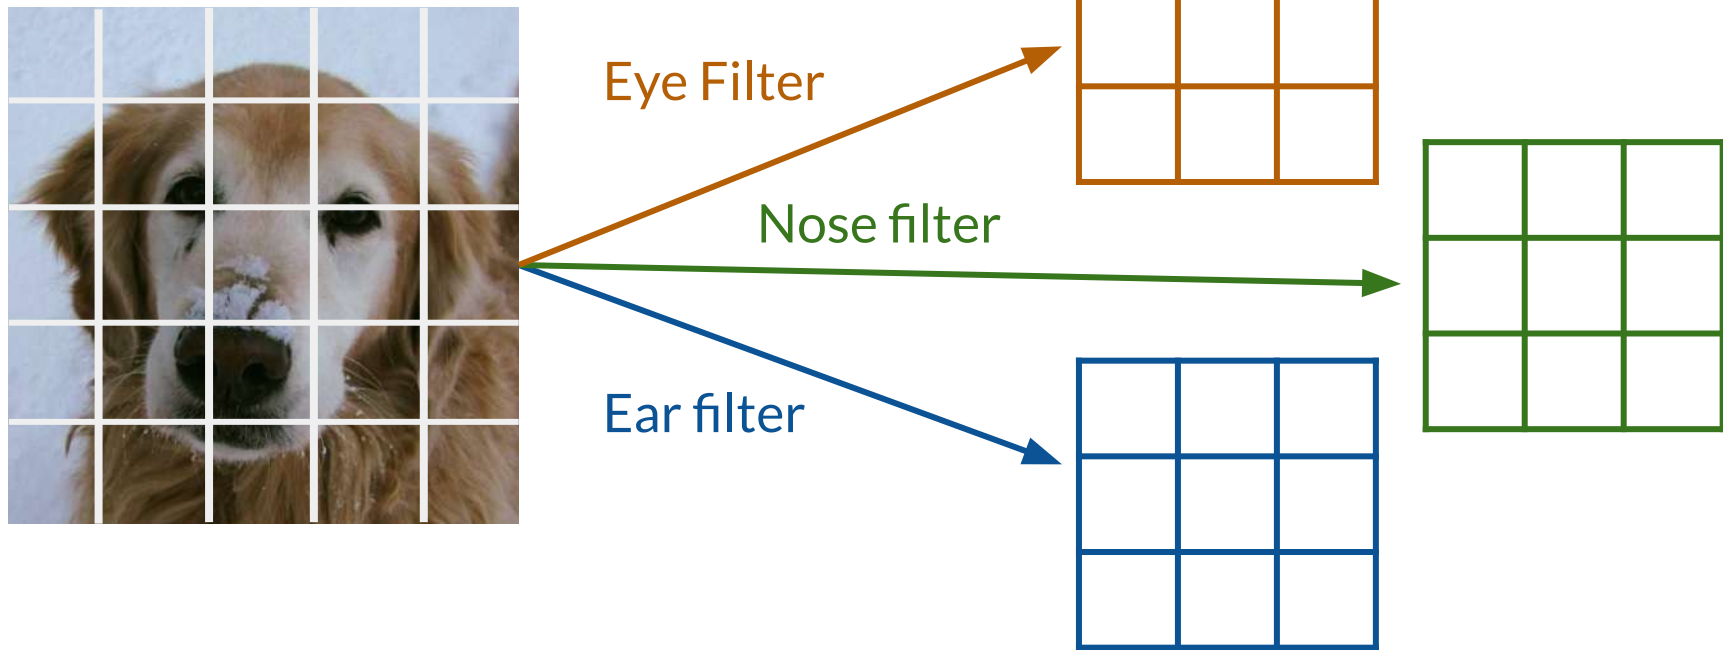
Eye Filter




What is a convolution?



What is a convolution?



What is a convolution?



What is a convolution?

50	50	0	0	0
50	50	0	0	0
50	50	0	0	0
50	50	0	0	0
50	50	0	0	0

Grayscale image

0 (black) to 255 (white)

What is a convolution?

50	50	0	0	0
50	50	0	0	0
50	50	0	0	0
50	50	0	0	0
50	50	0	0	0

Grayscale image

*

Filter

1	0	-1
1	0	-1
1	0	-1

0 (black) to 255 (white)

What is a convolution?

50	50	0	0	0
50	50	0	0	0
50	50	0	0	0
50	50	0	0	0
50	50	0	0	0

Grayscale image



Convolution

Filter

1	0	-1
1	0	-1
1	0	-1

0 (black) to 255 (white)

What is a convolution?

50×1	50×0	0×-1	0	0
50×1	50×0	0×-1	0	0
50×1	50×0	0×-1	0	0
50	50	0	0	0
50	50	0	0	0

Grayscale image




Convolution

Filter


1	0	-1
1	0	-1
1	0	-1

0 (black) to 255 (white)


What is a convolution?




What is a convolution?




What is a convolution?




What is a convolution?



What is a convolution?



What is a convolution?



Summary

- Convolutions are useful layers for processing images
- They scan the image to detect useful features
- Just element-wise products and sums!






deeplearning.ai


Padding and Stride

Outline

- Padding and stride
- The intuition behind padding



Stride



Grayscale image

Stride

50*1	50*0	0*-1	0	0
50*1	50*0	0*-1	0	0
50*1	50*0	0*-1	0	0
50	50	0	0	0
50	50	0	0	0

Grayscale image

*

Filter

1	0	-1
1	0	-1
1	0	-1

=

Result

150		

Stride

→ 1 Pixel to the right

50	50*1	0*0	0*-1	0
50	50*1	0*0	0*-1	0
50	50*1	0*0	0*-1	0
50	50	0	0	0
50	50	0	0	0

Grayscale image

*

Filter

1	0	-1
1	0	-1
1	0	-1

=

Result

150	150	

Stride

→ 1 Pixel to the right

50	50	0*1	0*0	0*-1
50	50	0*1	0*0	0*-1
50	50	0*1	0*0	0*-1
50	50	0	0	0
50	50	0	0	0

Grayscale image

*

Filter

1	0	-1
1	0	-1
1	0	-1

=

Result

150	150	0

Stride

→ 1 Pixel to the right

↓ 1 Pixel down

50	50	0	0	0
50*1	50*0	0*-1	0	0
50*1	50*0	0*-1	0	0
50*1	50*0	0*-1	0	0
50	50	0	0	0

Grayscale image

*

Filter

1	0	-1
1	0	-1
1	0	-1

=

Result

150	150	0
150		

Stride

→ 1 Pixel to the right



1 Pixel down

50	50	0	0	0
50	50*1	0*0	0*-1	0
50	50*1	0*0	0*-1	0
50	50*1	0*0	0*-1	0
50	50	0	0	0

Grayscale image



Filter

1	0	-1
1	0	-1
1	0	-1



Result

150	150	0
150	150	

Stride

→ 1 Pixel to the right



1 Pixel down

50	50	0	0	0
50	50	0	0	0
50	50	0	0	0
50	50	0	0	0
50	50	0	0	0

Grayscale image



Filter

1	0	-1
1	0	-1
1	0	-1

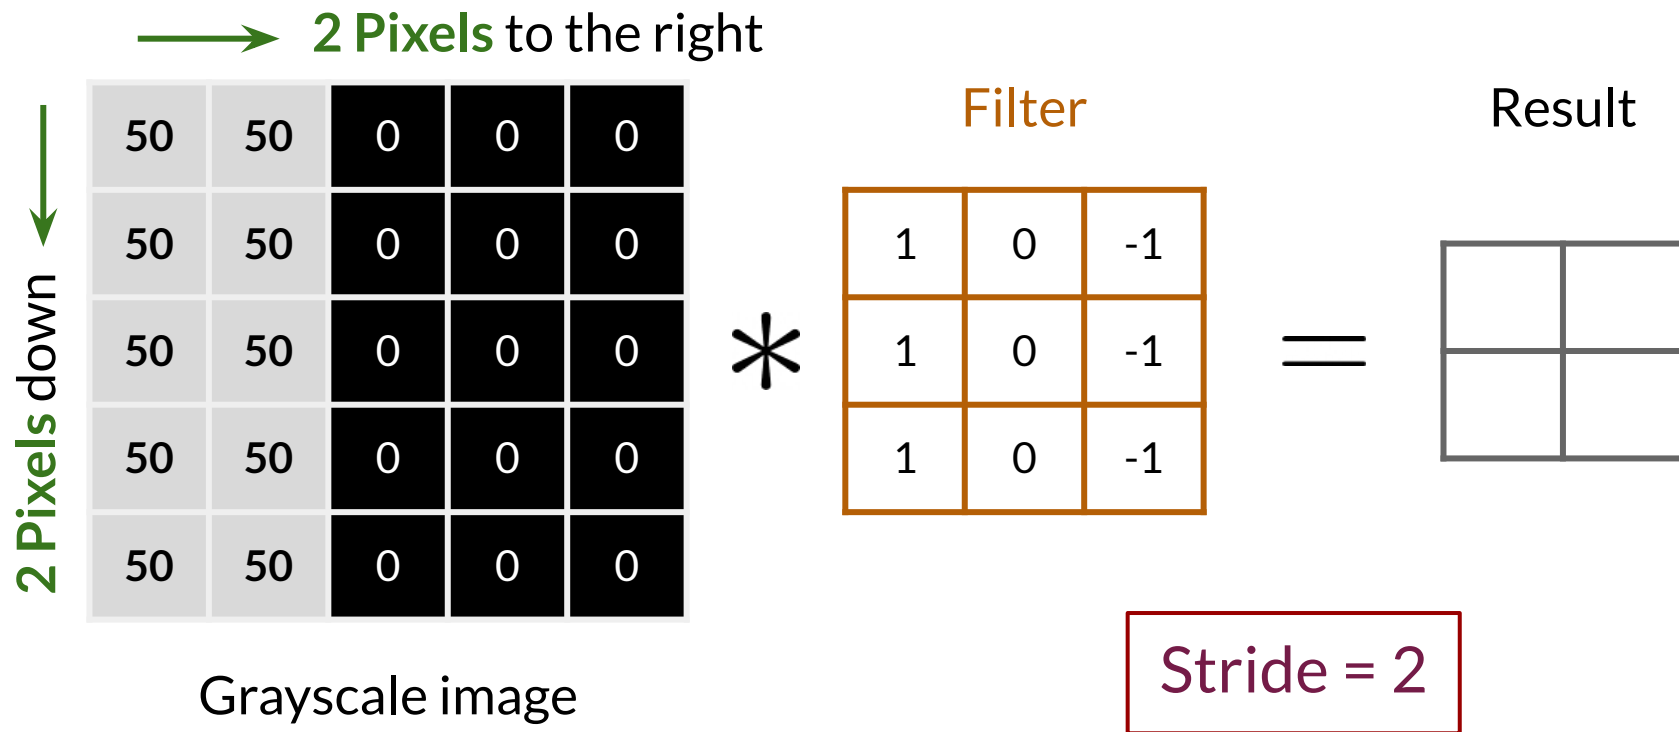


Result

150	150	0
150	150	0
150	150	0

Stride = 1

Stride



Stride

→ 2 Pixels to the right

↓ 2 Pixels down

50*1	50*0	0*-1	0	0
50*1	50*0	0*-1	0	0
50*1	50*0	0*-1	0	0
50	50	0	0	0
50	50	0	0	0

Grayscale image

*

Filter

1	0	-1
1	0	-1
1	0	-1


=

Result

150	

Stride = 2

Stride



Stride

→ 2 Pixels to the right

↓ 2 Pixels down

50	50	0	0	0
50	50	0	0	0
50*1	50*0	0*-1	0	0
50*1	50*0	0*-1	0	0
50*1	50*0	0*-1	0	0

Grayscale image

*

Filter

1	0	-1
1	0	-1
1	0	-1

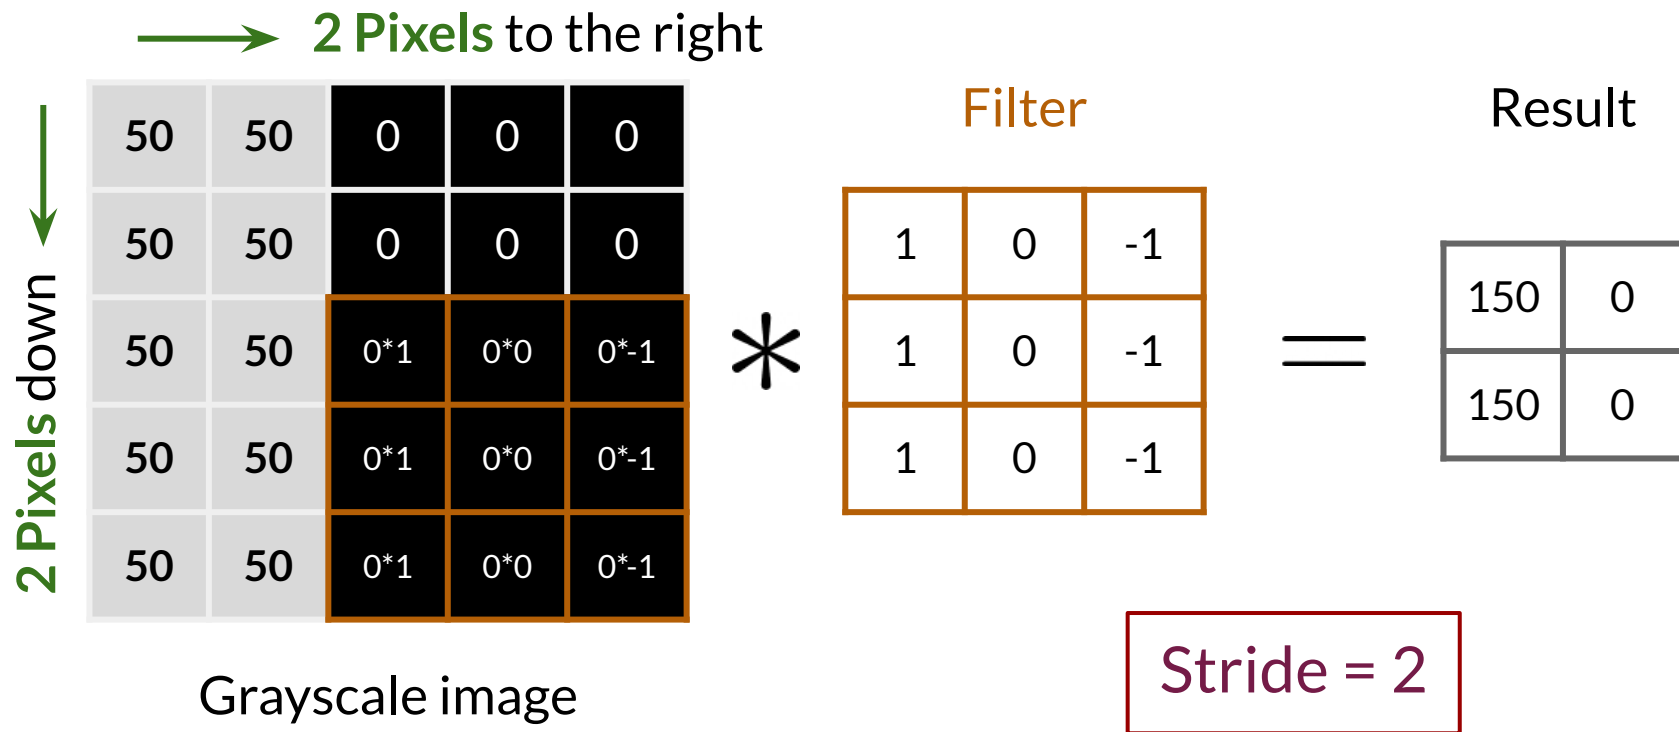
=

Result


150	0
150	

Stride = 2


Stride




Stride



Padding




Padding



Summary

- Stride determines how the filter scans the image
- Padding is like a frame on the image
- Padding gives similar importance to the edges and the center






deeplearning.ai


Pooling and Upsampling

Outline


- Pooling
- Upsampling and its relation to pooling



Pooling



Max Pooling



4x4 input to 2x2 output

Max Pooling

0	20	30	45
8	1	23	43
0	4	40	20
10	6	43	42

2x2 pooling with stride = 2

Max pooling

4x4 input to 2x2 output

Max Pooling

0	20	30	45
8	1	23	43
0	4	40	20
10	6	43	42


2x2 pooling with stride = 2

Max pooling

20	


4x4 input to 2x2 output

Max Pooling




4x4 input to 2x2 output

Max Pooling



4x4 input to 2x2 output

Max Pooling



Max Pooling

0	20	30	45
8	1	23	43
0	4	40	20
10	6	43	42

2x2 pooling with stride = 2

Max pooling


20	45
10	43

4x4 input to 2x2 output

Other types include:

1. Average pooling
2. Min pooling

Upsampling



Upsampling



Upsampling: Nearest Neighbors


2x2 input to 4x4 output

20	45
10	43

Upsampling
→


Upsampling: Nearest Neighbors

2x2 input to 4x4 output



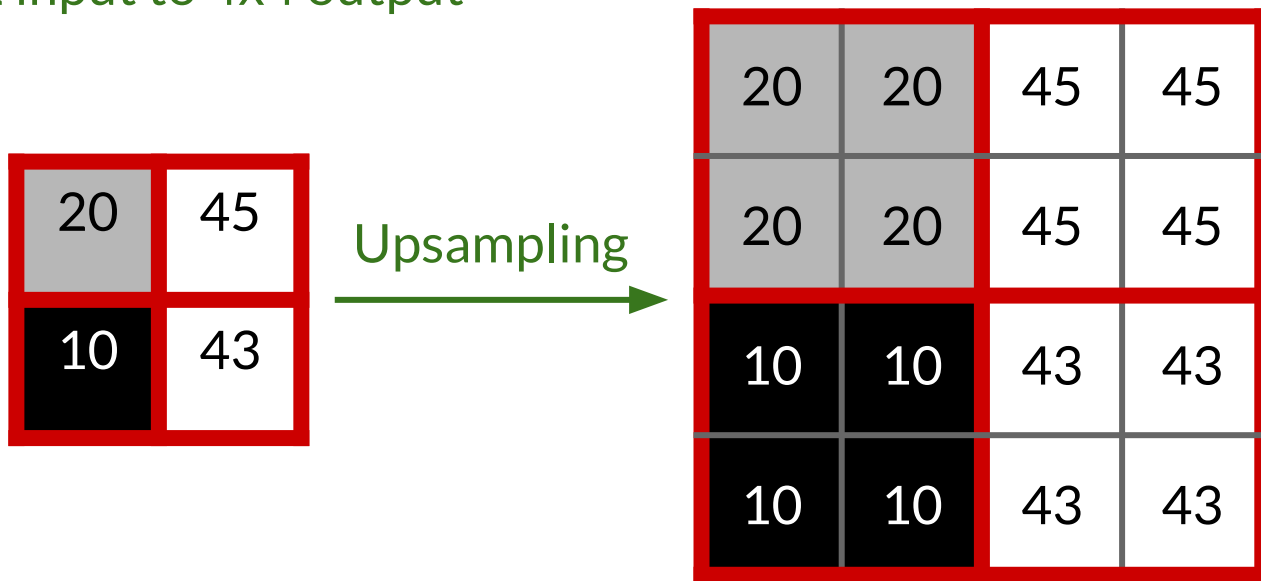
Upsampling: Nearest Neighbors

2x2 input to 4x4 output



Upsampling: Nearest Neighbors

2x2 input to 4x4 output



Upsampling: Nearest Neighbors

2x2 input to 4x4 output

20	45
10	43

Upsampling →


20	20	45	45
20	20	45	45
10	10	43	43
10	10	43	43

Other types include:

1. Linear interpolation
2. Bi-linear interpolation

Summary

- Pooling reduces the size of the input
- Upsampling increases the size of the input
- No learnable parameters!






deeplearning.ai

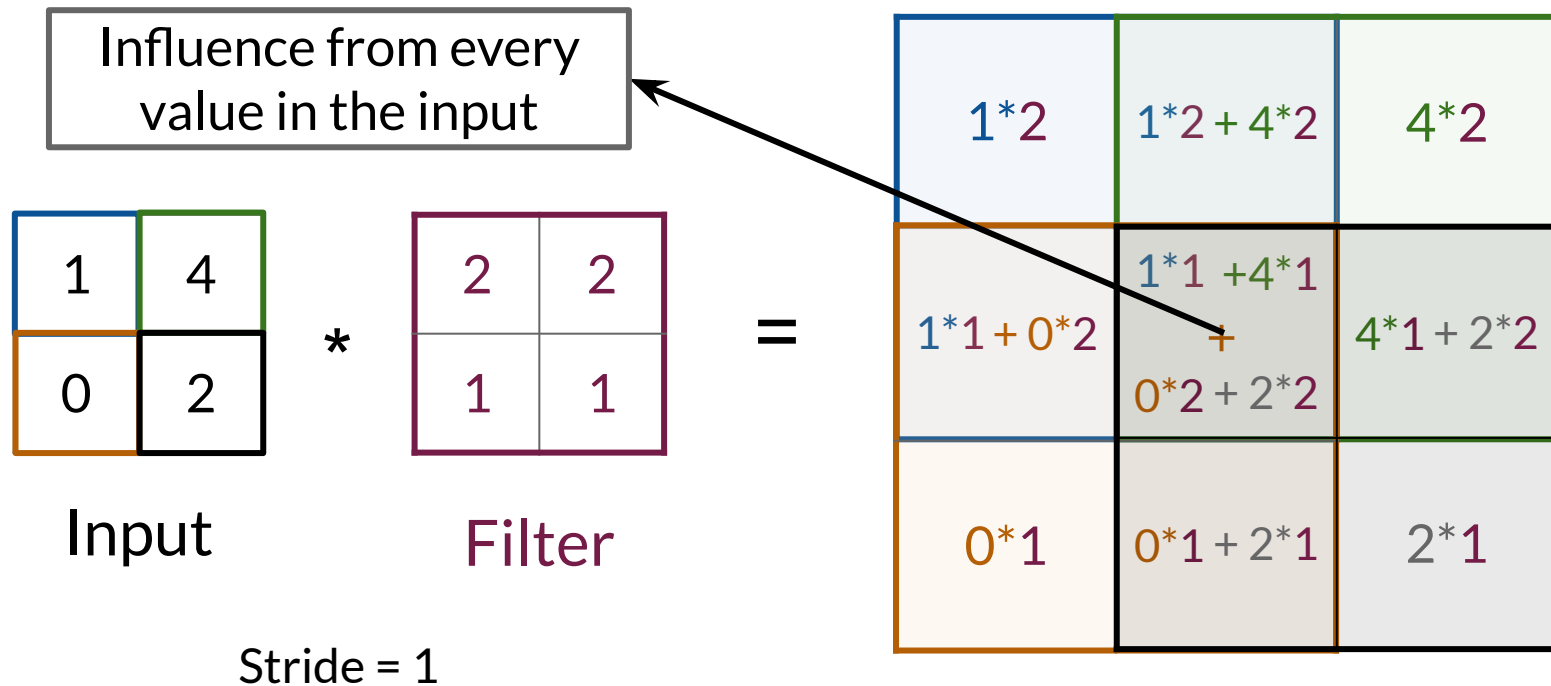
Transposed Convolutions

Outline


- Transposed convolutions as an upsampling technique
- Issues with transposed convolutions




Transposed Convolution



Transposed Convolution



The Problems with Transposed Convolution




Checkerboard
Pattern

Available from: <http://doi.org/10.23915/distill.00003>

Summary

- Transposed convolutions upsample
- They have learnable parameters
- Problem: results have a checkerboard pattern



Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>




deeplearning.ai


Mode Collapse

Outline


- Modes in distributions
- Mode collapse in GANs
- Intuition behind it during training



Mode Collapse




Mode Collapse



Mode Collapse



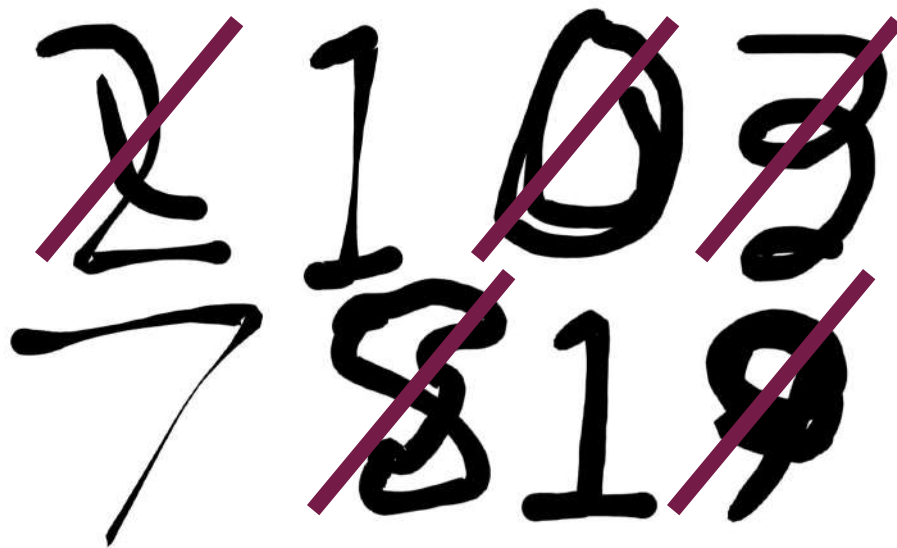
Discriminator



Mode Collapse



Discriminator




Fakes

Mode Collapse




Generator



Mode Collapse



Generator




Fakes that
fooled the
discriminator

Mode Collapse




Generator



Mode Collapse



Discriminator




Fakes

Mode Collapse




Generator



Mode Collapse



Generator




Fakes that
fooled the
discriminator

Mode Collapse



Generator



Summary

- Modes are peaks in the distribution of features
- Typical with real-world datasets
- Mode collapse happens when the generator gets stuck in one mode






deeplearning.ai

Problem with BCE Loss

Outline

- BCE Loss and the end objective in GANs
- Problem with BCE Loss



BCE Loss in GANs

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

Prediction

Label

Features

Parameters



Generator


Maximize
cost



Discriminator


Minimize
cost

Objective in GANs



Objective in GANs

Make the generated and real distributions look similar



BCE Loss in GANs


Criticizing is more straightforward



Discriminator

Single output

Easier to train
than the
generator




Generator

Complex
output


Difficult to
train

Often, the discriminator gets better than the generator



Problems with BCE Loss



Problems with BCE Loss




Problems with BCE Loss



Summary

- GANs try to make the real and generated distributions look similar
- When the discriminator improves too much, the function approximated by BCE Loss will contain flat regions
- Flat regions on the cost function = **vanishing gradients**






deeplearning.ai


Earth Mover's Distance

Outline

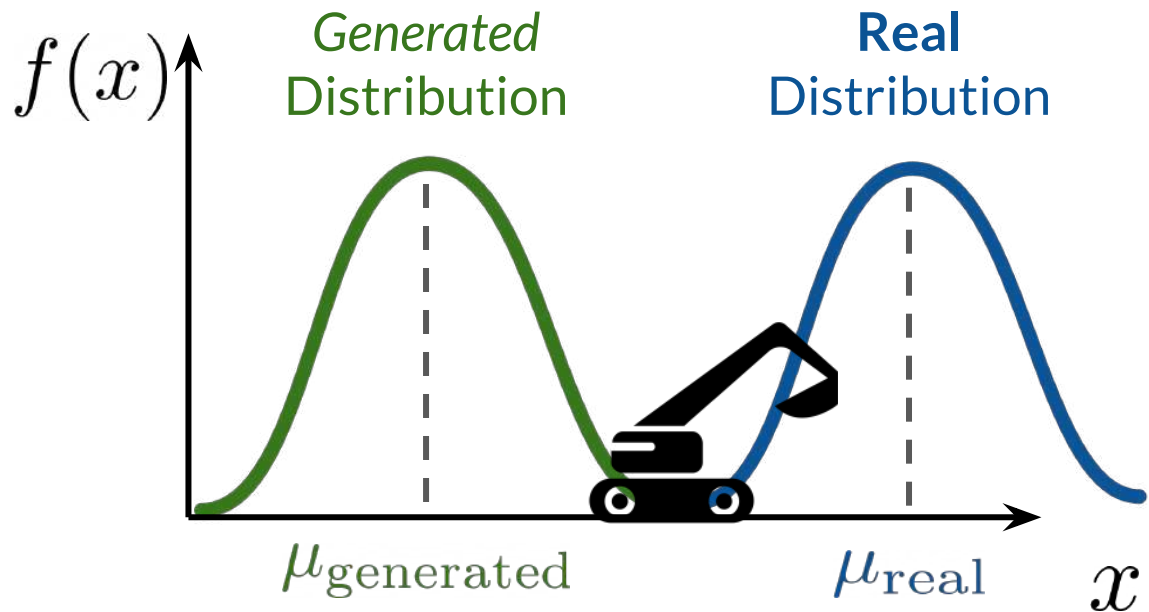
- Earth Mover's Distance (EMD)
- Why it solves the vanishing gradient problem of BCE Loss



Earth Mover's Distance





Earth Mover's Distance



Effort to make the *generated* distribution equal to the *real* distribution

Depends on the distance and amount moved

Earth Mover's Distance



Summary

- Earth mover's distance (EMD) is a function of amount and distance
- Doesn't have flat regions when the distributions are very different
- Approximating EMD solves the problems associated with BCE





deeplearning.ai

Wasserstein Loss

Outline

- BCE Loss Simplified
- W-Loss and its comparison with BCE Loss



BCE Loss Simplified

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

$\min_d \max_g$



Discriminator

Minimize
cost



Generator

Maximize
cost

BCE Loss Simplified

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

$$\min_d \max_g -[\mathbb{E}(\log(d(x))) + \mathbb{E}(\log(g(x)))]$$



Minimize
cost



Maximize
cost

BCE Loss Simplified

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

$$\min_d \max_g -[\mathbb{E}(\log(d(x))) + \mathbb{E}(1 - \log(d(g(z))))]$$



Minimize
cost



Maximize
cost

W-Loss

W-Loss approximates the Earth Mover's Distance

W-Loss

W-Loss approximates the Earth Mover's Distance

$$\min_g \max_c \mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))$$

W-Loss

W-Loss approximates the Earth Mover's Distance

$$\min_g \max_c \mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))$$



Maximize
the
distance

W-Loss

W-Loss approximates the Earth Mover's Distance

$$\min_g \max_c \mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))$$



Generator

Minimize
the
distance




Critic


Maximize
the
distance

Discriminator Output

Discriminator output





Discriminator output



Discriminator Output

Discriminator output

~~Discriminator output~~
Critic



W-Loss vs BCE Loss

BCE Loss

W-Loss

Discriminator outputs between 0 and 1

$$-[\mathbb{E}(\log(d(x))) + \mathbb{E}(1 - \log(d(g(z))))]$$


Critic outputs any number

$$\mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))$$

W-Loss helps with mode collapse and vanishing gradient problems

Summary

- W-Loss looks very similar to BCE Loss
- W-Loss prevents mode collapse and vanishing gradient problems





deeplearning.ai

Condition on Wasserstein Critic

Outline

- Continuity condition on the critic's neural network
- Why this condition matters



Condition on W-Loss

$$\min_g \max_c \mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))$$

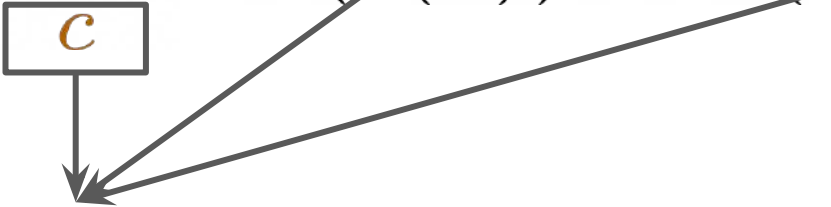
Condition on W-Loss

$$\min_g \max_c \mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))$$

Condition on W-Loss

$$\min_g \max_c \mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))$$

Condition on W-Loss


$$\min_g \max_c \mathbb{E}(\boxed{c}(x)) - \mathbb{E}(\boxed{c}(g(z)))$$
A diagram with two arrows pointing from the boxed 'c' in the equation to the text 'Needs to be 1-Lipschitz Continuous'. One arrow starts from the bottom of the box and points down. The other starts from the top-right corner of the box and points diagonally down and to the right.

Needs to be 1-Lipschitz Continuous

Condition on W-Loss

Critic needs to be **1-L** Continuous


The norm of the gradient should be at most **1** for *every point*



Condition on W-Loss


Critic needs to be **1-L** Continuous

The norm of the gradient should be at most **1** for every point



Condition on W-Loss

Critic needs to be **1-L** Continuous




The norm of the gradient should be at most **1** for every point

Condition on W-Loss

Critic needs to be **1-L Continuous**

The norm of the gradient should be at most **1** for every point




Not 1-L Continuous

Condition on W-Loss

Critic needs to be **1-L** Continuous


The norm of the gradient should be at most **1** for *every point*



Condition on W-Loss

Critic needs to be **1-L** Continuous


The norm of the gradient should be at most **1** for every point



Condition on W-Loss

Critic needs to be **1-L** Continuous


The norm of the gradient should be at most **1** for every point



Condition on W-Loss

Critic needs to be 1-L Continuous


The norm of the gradient should be at most **1** for every point



Condition on W-Loss

Critic needs to be 1-L Continuous

The norm of the gradient should be at most **1** for every point




1-L Continuous

W-Loss is valid

Needed for training stable
neural networks with W-Loss

Summary

- Critic's neural network needs to be 1-L Continuous when using W-Loss
- This condition ensures that W-Loss is validly approximating Earth Mover's Distance





deeplearning.ai

1-Lipschitz Continuity Enforcement


Outline

- Weight clipping and gradient penalty
- Advantages of gradient penalty



1-L Enforcement

Critic needs to be 1-L Continuous




Norm of the gradient at most 1

$$||\nabla f(x)||_2 \leq 1$$

Slope of the function
at most 1

1-L Enforcement: Weight Clipping

Weight clipping forces the weights of the critic to a fixed interval



Limits the learning ability of the critic

1-L Enforcement: Gradient Penalty

$$\min_g \max_c \mathbb{E}(c(x)) - \mathbb{E}(c(g(z))) + \lambda \text{reg}$$



Regularization of the
critic's gradient

1-L Enforcement: Gradient Penalty

Real



ϵ

Random interpolation



1-L Enforcement: Gradient Penalty

Real



Random interpolation



ϵ

$1 - \epsilon$

\hat{x}

Generated



1-L Enforcement: Gradient Penalty

$$\mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2$$

Regularization term

1-L Enforcement: Gradient Penalty

$$\mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2$$

Regularization term

1-L Enforcement: Gradient Penalty

$$\mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2$$

Regularization term



$$\epsilon x + (1 - \epsilon)g(z)$$

Interpolation

1-L Enforcement: Gradient Penalty

$$\mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2$$

Regularization term



$$\epsilon \boxed{x} + (1 - \epsilon)g(z)$$

Real

Interpolation

1-L Enforcement: Gradient Penalty

$$\mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2$$

Regularization term



$$\epsilon \boxed{x} + (1 - \epsilon) \boxed{g(z)}$$

Interpolation

Real

Generated

Putting It All Together

$$\min_g \max_c \mathbb{E}(c(x)) - \mathbb{E}(c(g(z))) + \lambda \mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2$$

Putting It All Together

$$\min_g \max_c \boxed{\mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))} + \lambda \mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2$$

Makes the GAN less prone to **mode collapse** and **vanishing gradient**

Putting It All Together


$$\min_g \max_c \mathbb{E}(c(x)) - \mathbb{E}(c(g(z))) + \lambda \mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2$$

Makes the GAN less prone to **mode collapse** and **vanishing gradient**

Tries to make the critic be 1-L Continuous, for the loss function to be **continuous and differentiable**

Summary

- Weight clipping and gradient penalty are ways to enforce 1-L continuity
- Gradient penalty tends to work better



Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>




deeplearning.ai

Conditional Generation: Intuition

Outline

- Unconditional generation
- Conditional vs. unconditional generation



Unconditional Generation

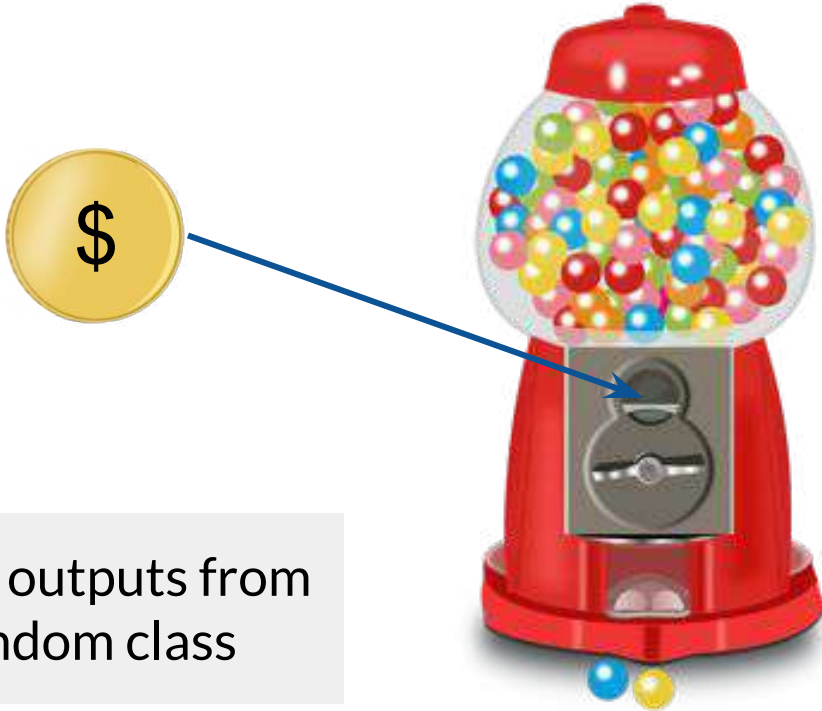
You get outputs from
a random class

Unconditional Generation

You get outputs from
a random class




Unconditional Generation

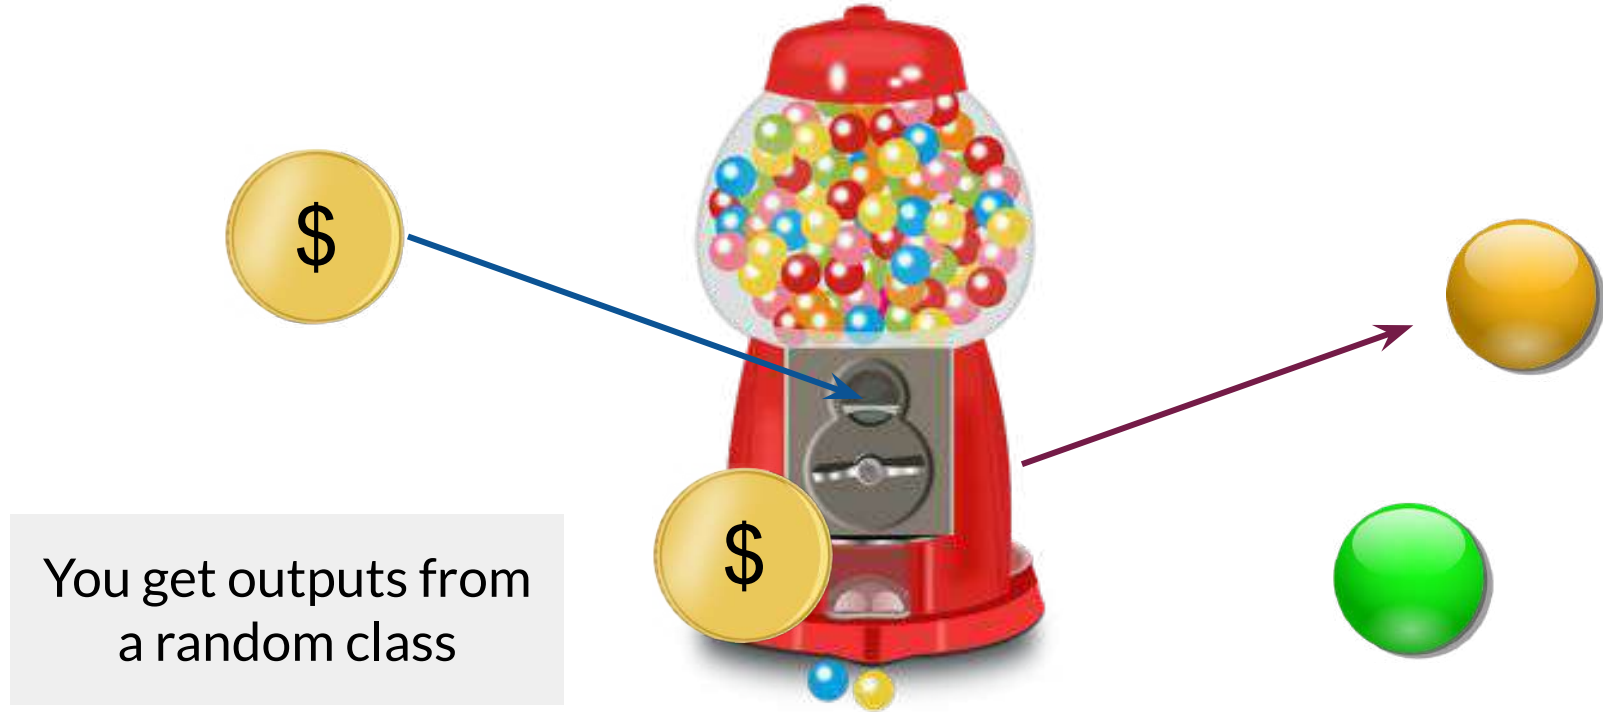


You get outputs from
a random class


Unconditional Generation



Unconditional Generation



Unconditional Generation

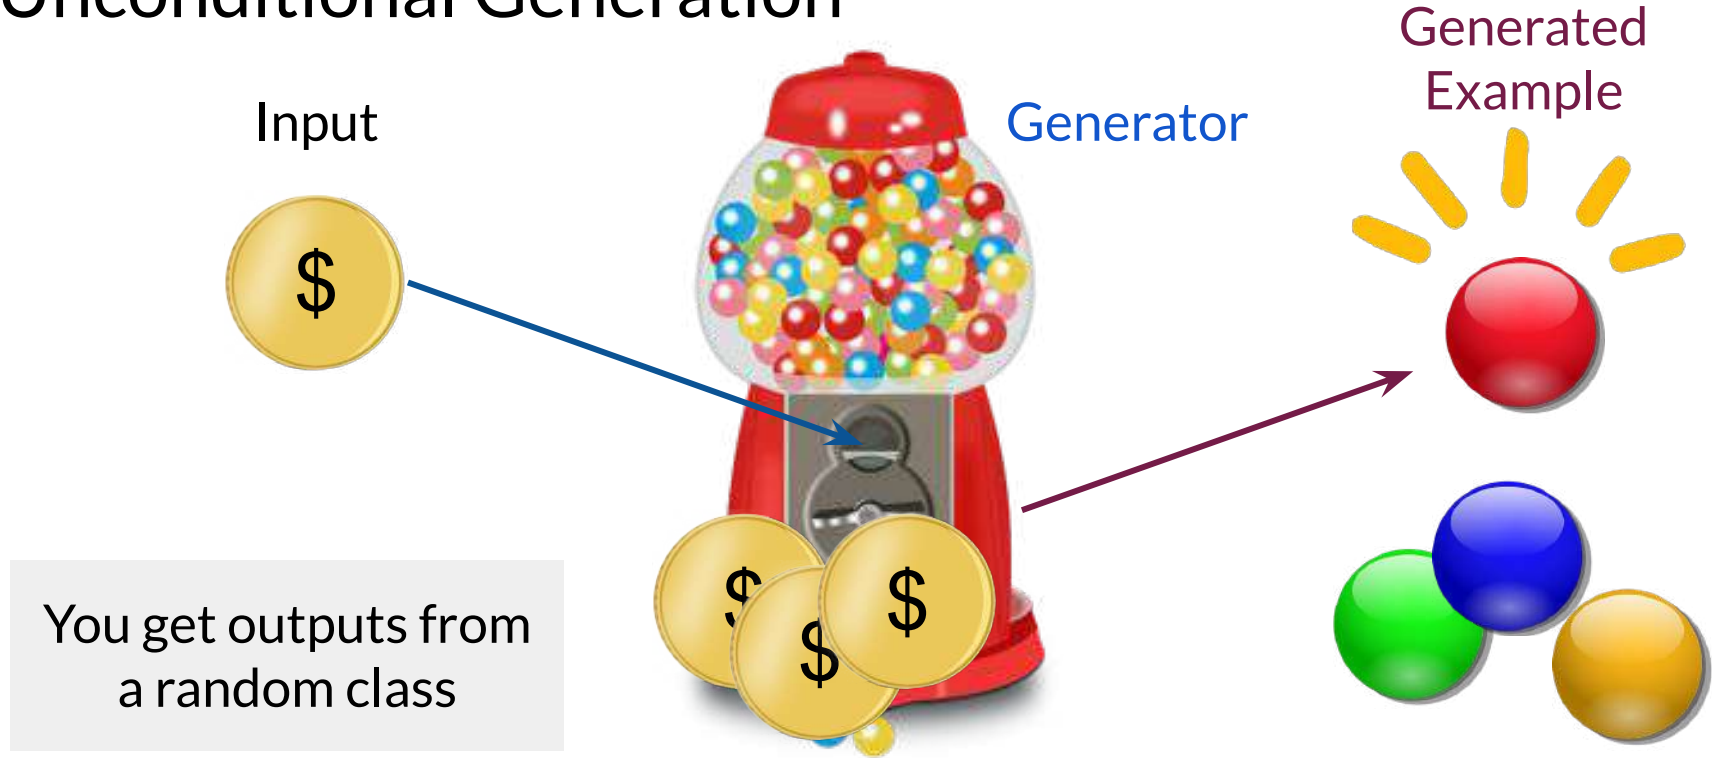


Unconditional Generation



You get outputs from
a random class


Unconditional Generation



Conditional Generation

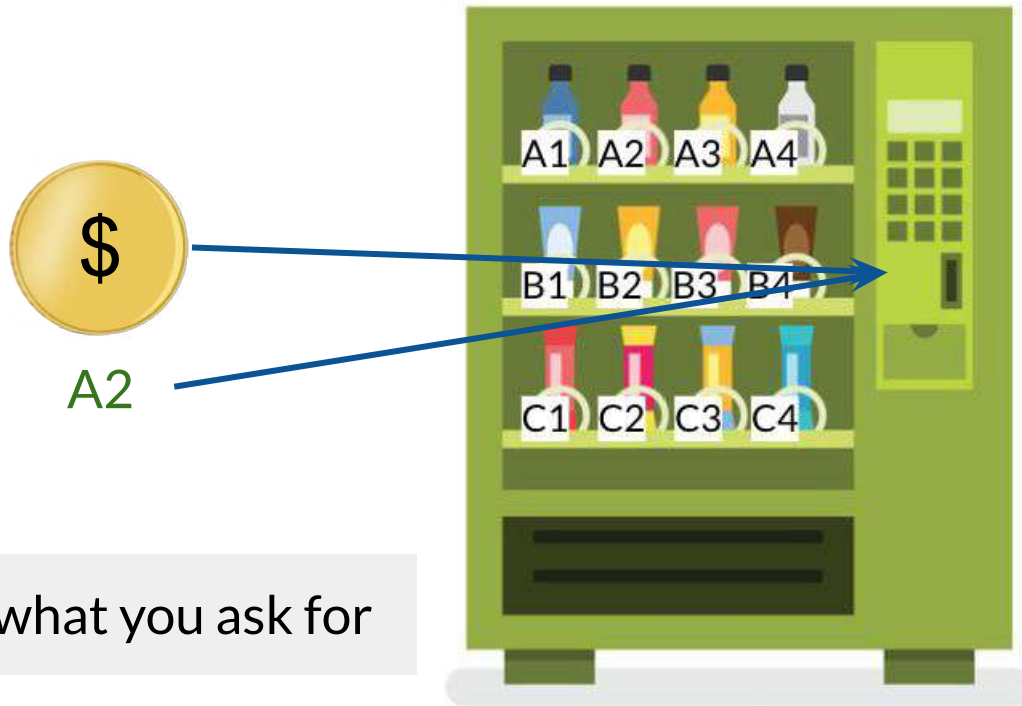
You get what you ask for

Conditional Generation




You get what you ask for

Conditional Generation




You get what you ask for

Conditional Generation



Conditional Generation



Conditional vs. Unconditional Generation



Conditional vs. Unconditional Generation

Conditional

Unconditional

Examples from **the classes
you want**


Examples from *random classes*

Conditional vs. Unconditional Generation

Conditional	Unconditional
Examples from the classes you want	Examples from <i>random classes</i>
Training dataset needs to be labeled	Training dataset <i>doesn't need to be labeled</i>

Summary

- Conditional generation requires labeled datasets
- Examples can be generated for the selected class






deeplearning.ai


Conditional Generation: Inputs

Outline

- How to tell the generator what type of example to produce
- Input representation for the discriminator




Generator Input




Generator Input

Noise vector




Class (one-hot) vector




Generator Input

Noise vector




Class (one-hot) vector



Husky


Generator Input

Noise vector



Randomness in the generation


Class (one-hot) vector



Husky


Generator Input

Noise vector



Randomness in the generation


Class (one-hot) vector




Husky

Control in the generation

Generator Input




Generator Input



Generator Input

Output




Discriminator Input



Discriminator

Discriminator Input



Discriminator Input




Golden Retriever



Output

5%
Real

Discriminator Input



Discriminator Input



Golden Retriever




Output

98%
Real

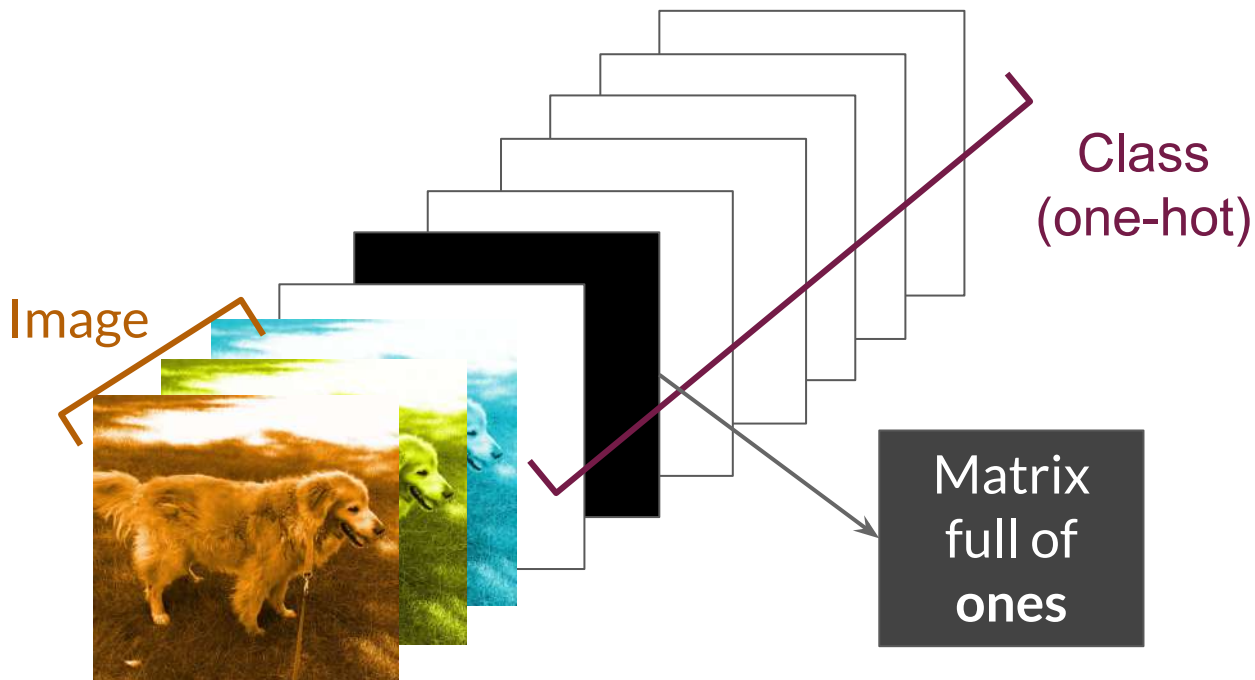
Discriminator Input




Discriminator Input



Discriminator Input



Discriminator Input



Summary

- The class is passed to the generator as one-hot vectors
- The class is passed to the discriminator as one-hot matrices
- The size of the vector and the number of matrices represent the number of classes






deeplearning.ai

Controllable Generation

Outline

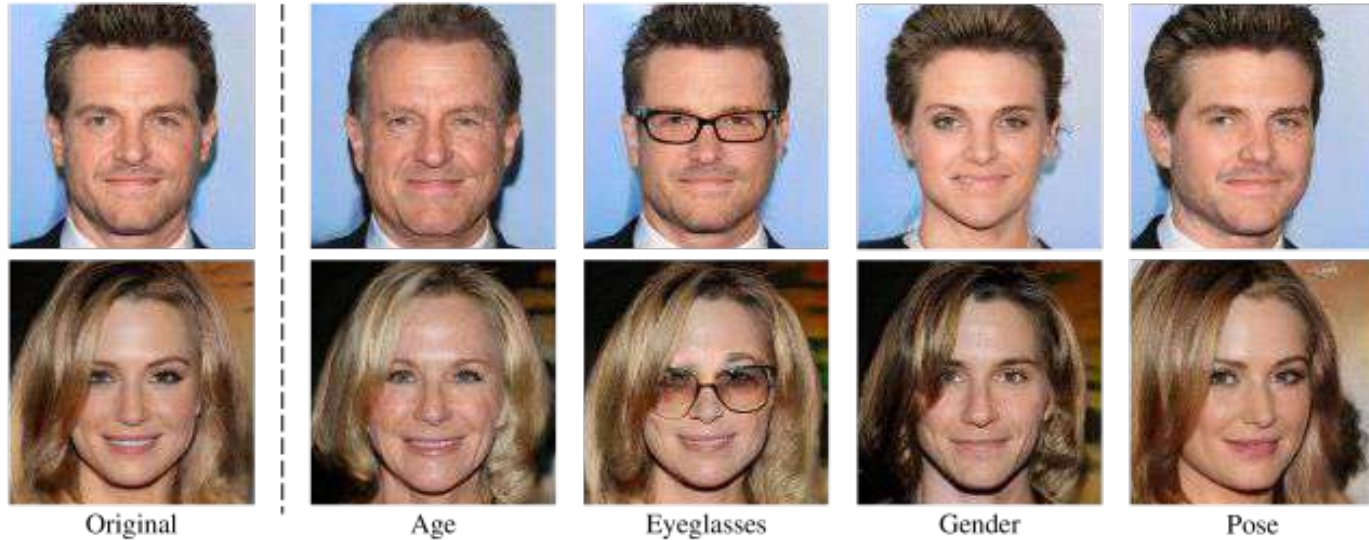
- What is controllable generation
- How it compares to conditional generation



Controllable Generation

Change specific features of the output


Controllable Generation



Change specific features of the output

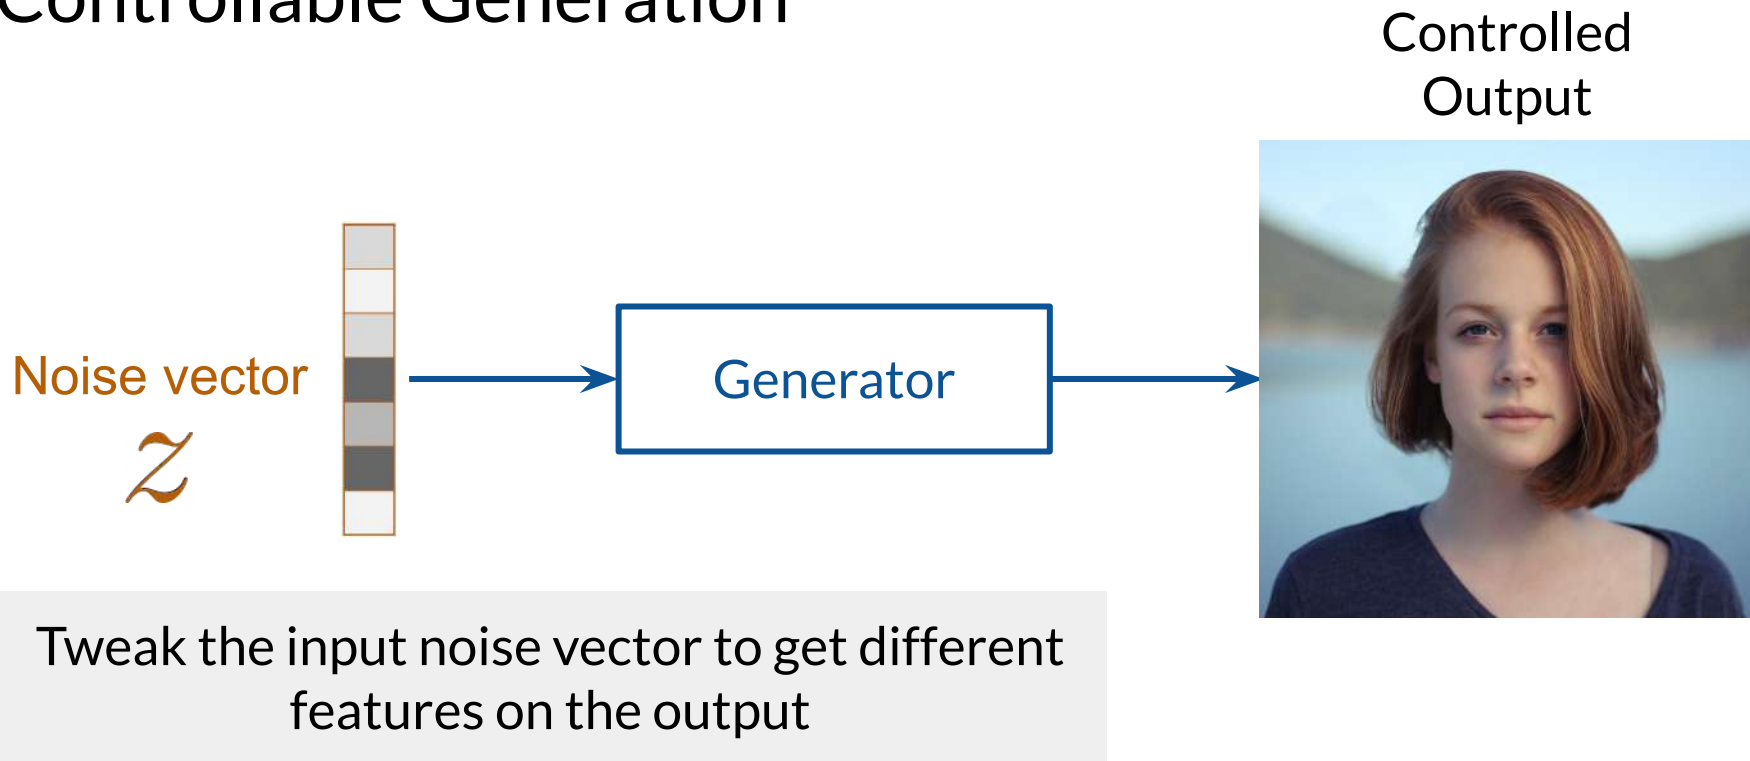
Available from: <https://arxiv.org/abs/1907.10786>

Controllable Generation

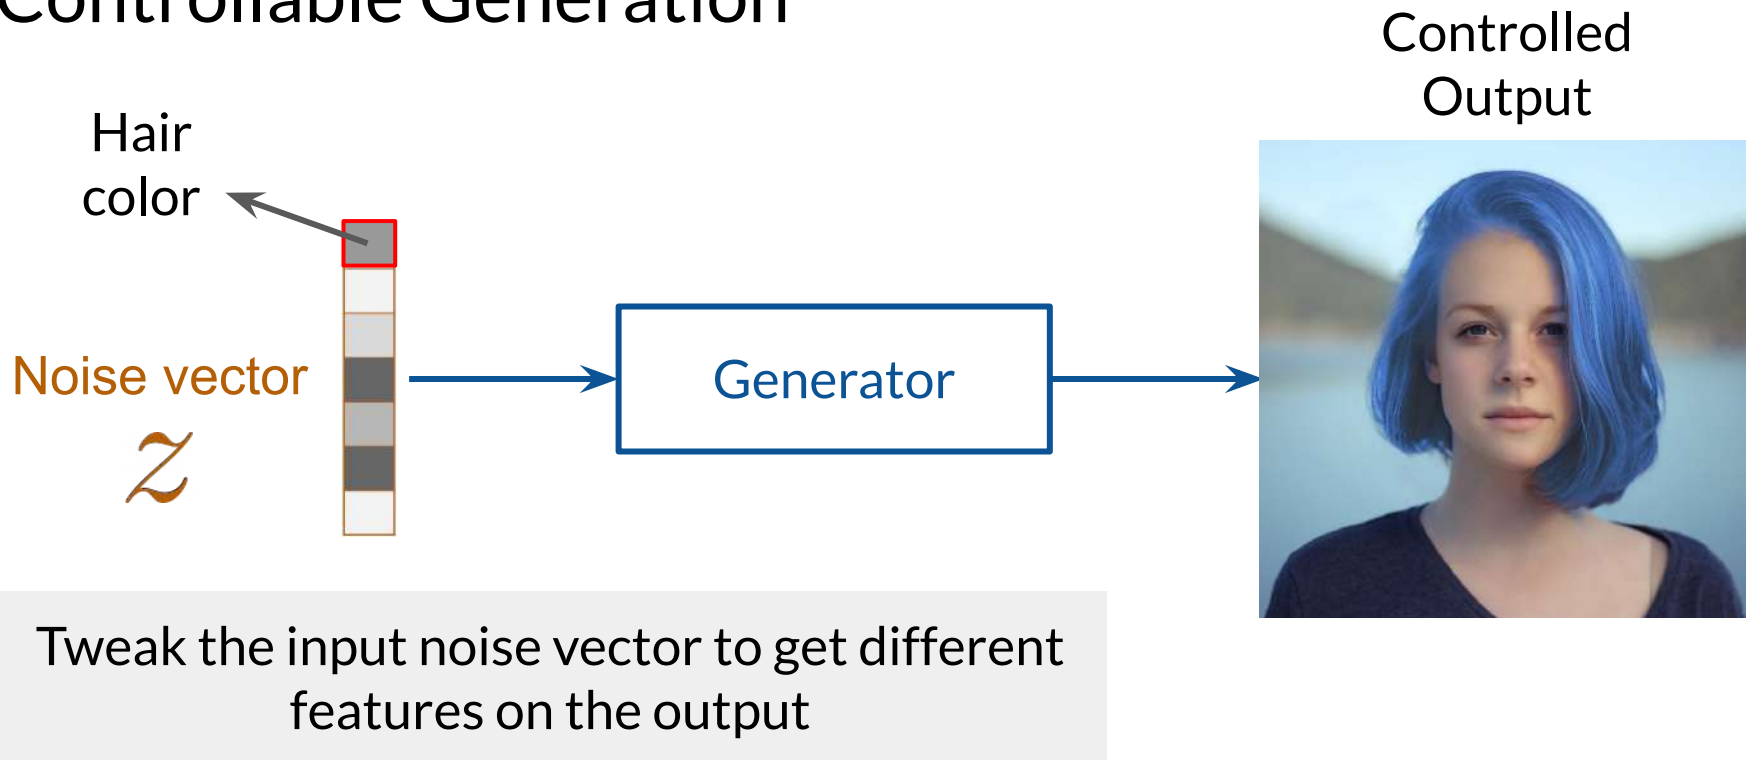


Tweak the input noise vector to get different features on the output

Controllable Generation



Controllable Generation



Controllable Generation vs. Conditional Generation

Controllable

Conditional

Controllable Generation vs. Conditional Generation

Controllable

Conditional

Examples with the **features**
that you want

Examples from *the classes you*
want

Controllable Generation vs. Conditional Generation

Controllable

Examples with the **features
that you want**

Training dataset **doesn't need
to be labeled**

Conditional

Examples from *the classes you
want*

Training dataset *needs to be
labeled*

Controllable Generation vs. Conditional Generation

Controllable

Examples with the **features that you want**

Training dataset **doesn't need to be labeled**

Manipulate the z vector
input

Conditional


Examples from *the classes you want*

Training dataset *needs to be labeled*

Append a class vector to the
input

Summary

- Controllable generation lets you control the features of the generated outputs
- It does not need a labeled training dataset
- The input vector is tweaked to get different features on the output






deeplearning.ai



Vector Algebra in the Z-Space

Outline

- Interpolation in the Z-space
- Modifying the noise vector z to control desired features




Interpolation Using the Z-Space




How an image morphs into another

Interpolation Using the Z-Space



Interpolation Using the Z-Space




$g(v_1)$



$g(v_2)$



Interpolation Using the Z-Space




$g(v_1)$



$g(v_2)$



Interpolation Using the Z-Space




Intermediate images using
the Z-space

$g(v_1)$ \longleftrightarrow $g(v_2)$




Z-Space and Controllable Generation





Generator

Z-Space and Controllable Generation




Output

Z-Space and Controllable Generation




Output


Z-Space and Controllable Generation



Z-Space and Controllable Generation



Z-Space and Controllable Generation




Original output

$$g(v_1) \longrightarrow$$



Z-Space and Controllable Generation




Original output

$$g(v_1) \longrightarrow$$



Z-Space and Controllable Generation



Original output

$$g(v_1) \rightarrow$$




Controlled output

$$g(v_1 + d) \rightarrow$$



Summary

- To control output features, you need to find directions in the Z-space
- To modify your output, you move around in the Z-space





deeplearning.ai

Challenges with Controllable Generation

Outline

- Output feature correlation
- Z-space entanglement



Feature Correlation

Uncorrelated
Features



Add beard



Feature Correlation

Uncorrelated
Features



Add beard




Correlated
Features



Add beard

Make more
masculine



Z-Space Entanglement

Noise vector

z



Glasses

Beard


Hair

Eyes


Age

Output
Features


Z-Space Entanglement




Z-Space Entanglement




Z-Space Entanglement



Z-Space Entanglement




Z-Space Entanglement



It is not possible to control single output features

Z-Space Entanglement



It is not possible to control single output features

Summary

- When trying to control one feature, others that are correlated change
- Z-space entanglement makes controllability difficult, if not impossible
- Entanglement happens when z does not have enough dimensions






deeplearning.ai

Classifier Gradients

Outline

- How to use classifiers to find directions in the Z-space
- Requirements to use this method




Classifier Gradients

Classifier Gradients

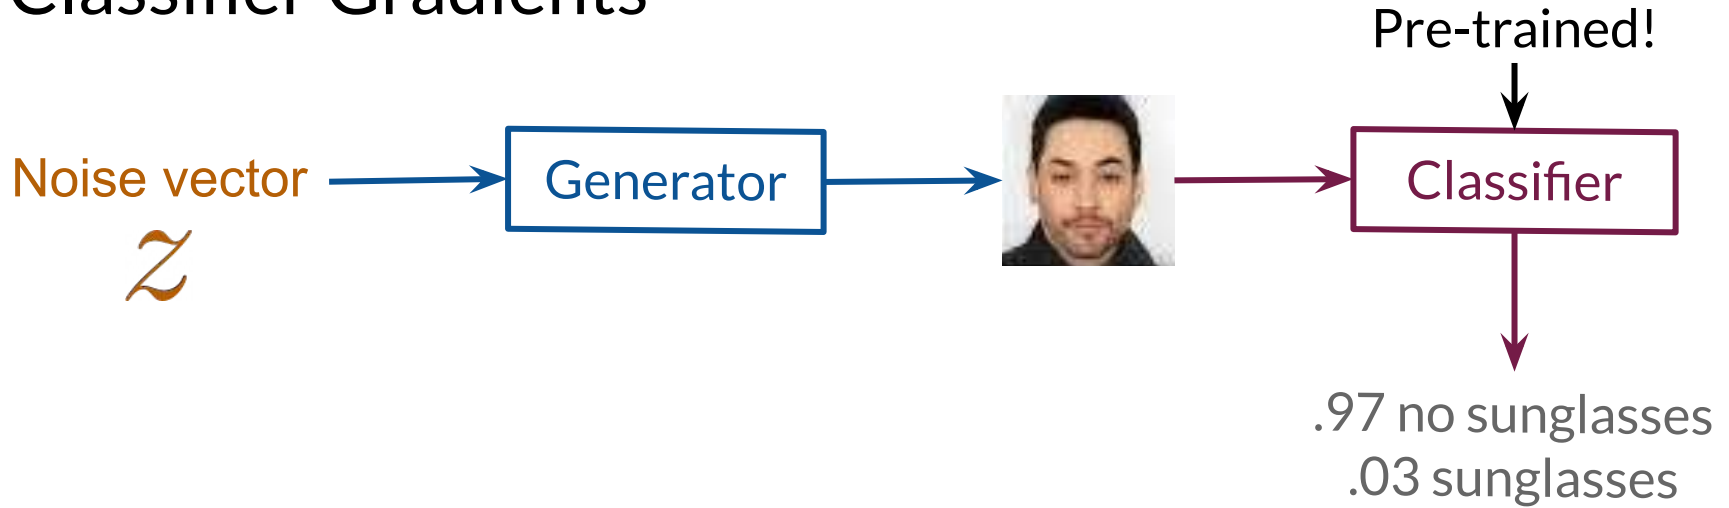
Noise vector

z

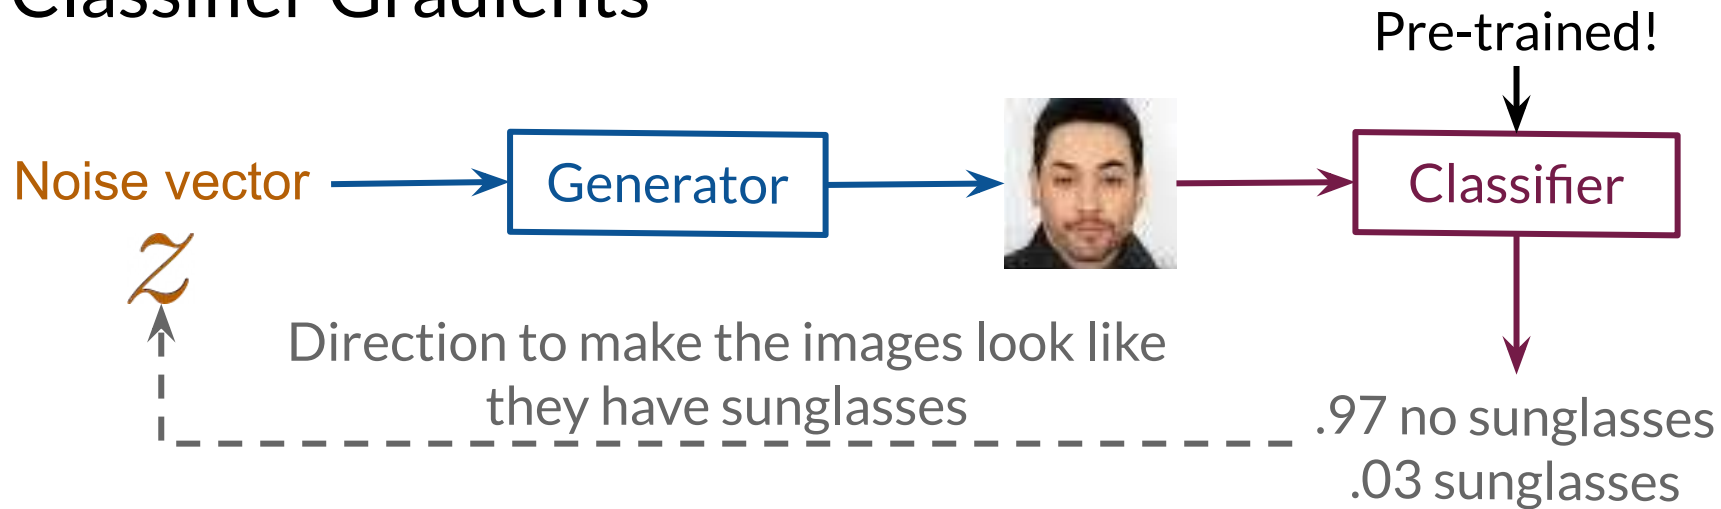
Classifier Gradients



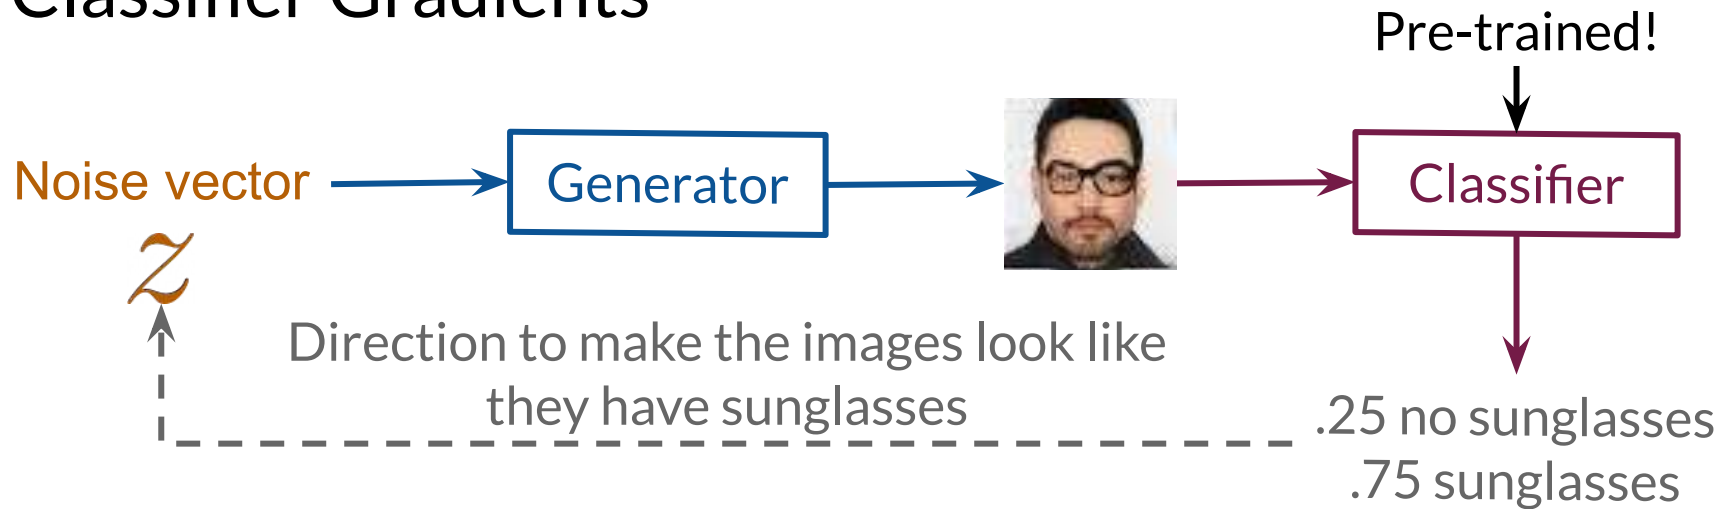
Classifier Gradients



Classifier Gradients

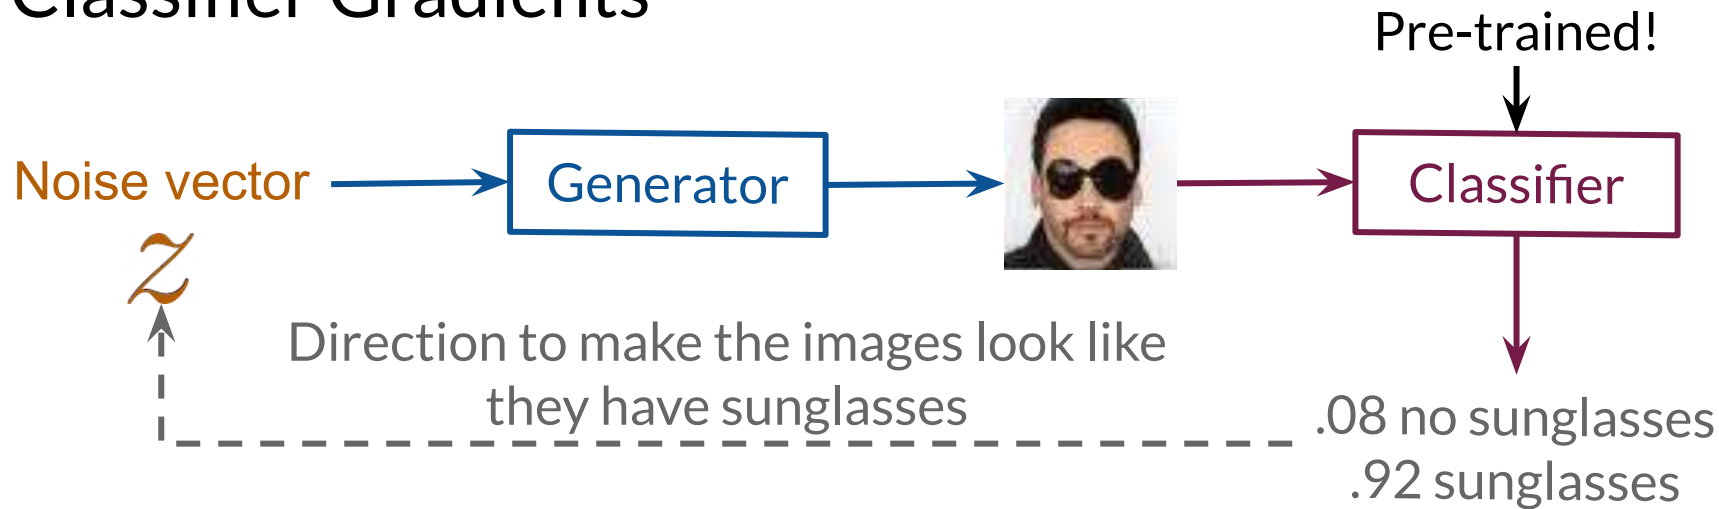


Classifier Gradients



Modify **just** the **noise vector** until the feature emerges

Classifier Gradients



Modify **just** the **noise vector** until the feature emerges

Summary

- Classifiers can be used to find directions in the Z-space
- To find directions, the updates are done just to the noise vector






deeplearning.ai

Disentanglement

Outline


- What a disentangled Z-space means
- Ways to encourage disentangled Z-spaces



Disentangled Z-Space

$$v_1 = [1, 2, 3, \dots]$$

$$v_2 = [5, 6, 7, \dots]$$




Disentangled Z-Space

$$v_1 = [\overset{z_1}{1}, 2, 3, \dots]$$


$$v_2 = [5, 6, 7, \dots]$$

Hair
color



Disentangled Z-Space


$$\begin{array}{c} z_1 \quad z_2 \\ v_1 = [\text{1}, \text{2}, 3, \dots] \\ v_2 = [\text{5}, \text{6}, 7, \dots] \\ \text{Hair color} \quad \text{Hair length} \end{array}$$



Disentangled Z-Space

$$\begin{array}{c} z_1 \quad z_2 \\ v_1 = [\text{1}, \text{2}, 3, \dots] \\ v_2 = [\text{5}, \text{6}, 7, \dots] \end{array}$$


Hair color Hair length



Disentangled Z-Space

$$\begin{array}{c} z_1 \quad z_2 \\ v_1 = [\text{1}, \text{2}, 3, \dots] \\ v_2 = [\text{5}, \text{6}, 7, \dots] \end{array}$$

Hair color Hair length




Disentangled Z-Space


$$\begin{array}{c} z_1 \quad z_2 \\ v_1 = [\text{1}, \text{2}, 3, \dots] \\ v_2 = [\text{5}, \text{6}, 7, \dots] \end{array}$$

Hair color Hair length


Latent factors of variation



Disentangled Z-Space




Disentangled Z-Space



Changes to one feature
don't affect the others


Disentangled Z-Space



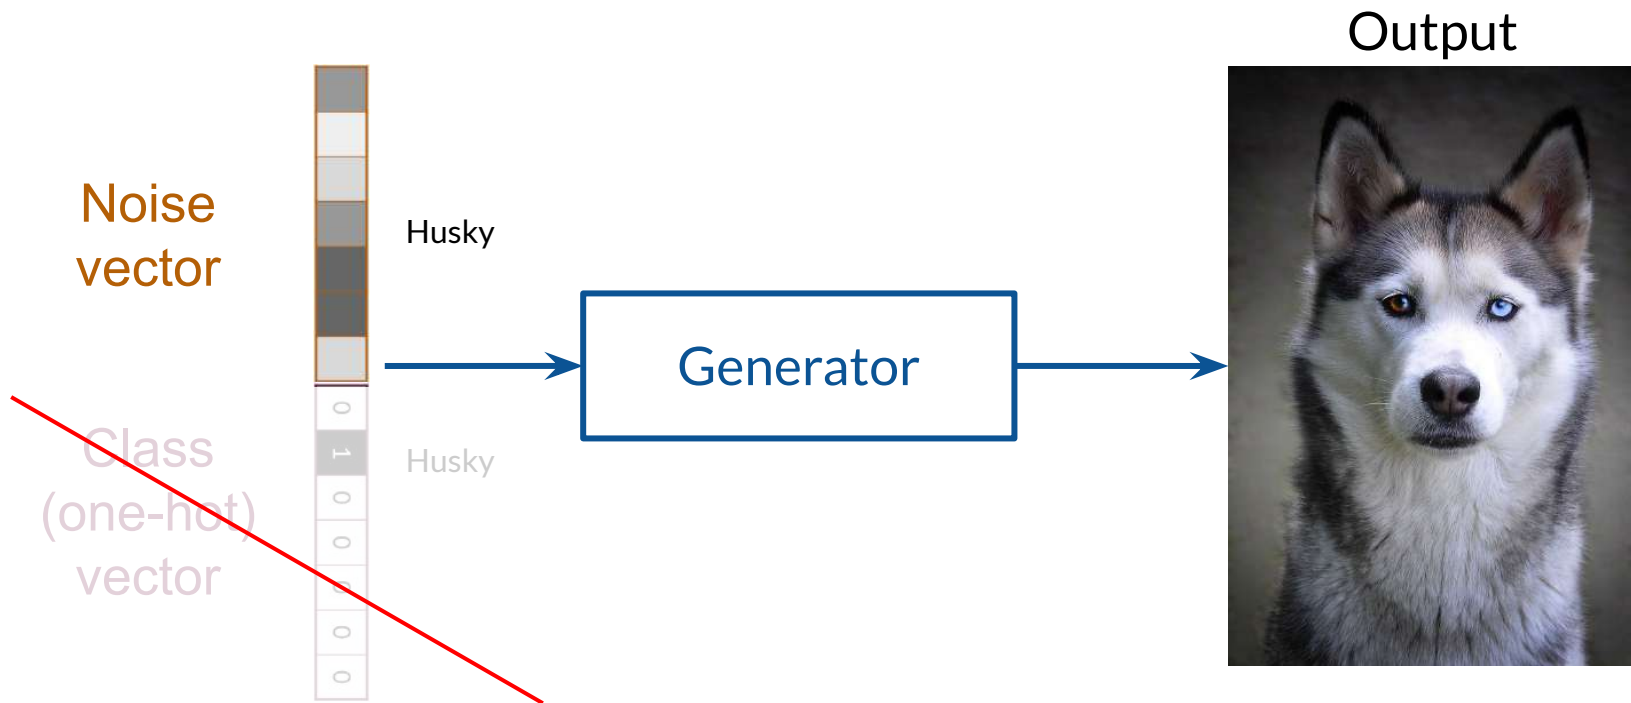
Changes to one feature
don't affect the others



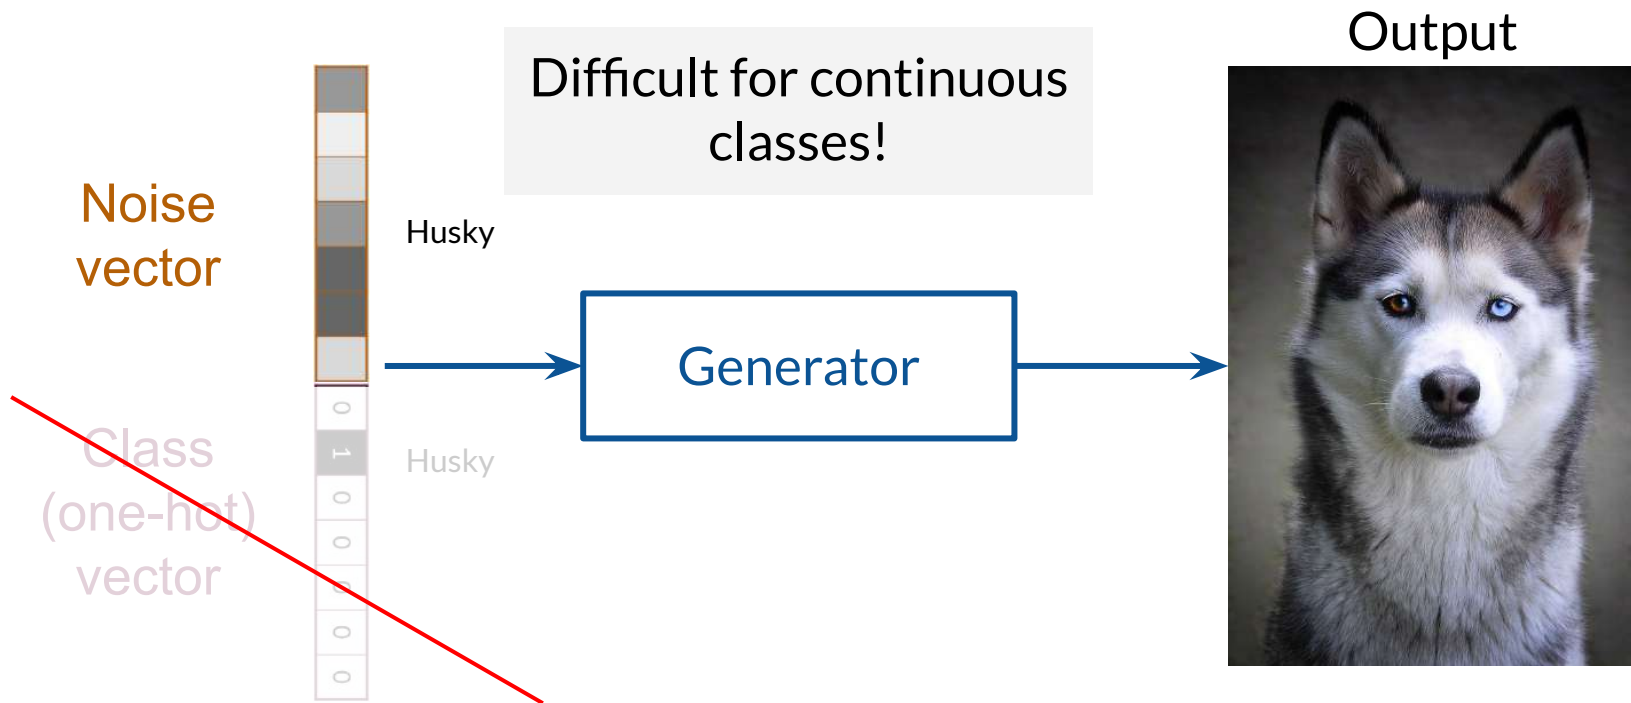
Encourage Disentanglement: Supervision



Encourage Disentanglement: Supervision



Encourage Disentanglement: Supervision



Encourage Disentanglement: Loss Function

$$v_1 = [1, 2, 3, \dots]$$

$$v_2 = [5, 6, 7, \dots]$$

Encourage Disentanglement: Loss Function

$$v_1 = [1, 2, 3, \dots]$$

$$v_2 = [5, 6, 7, \dots]$$

$$L_{\text{new}} = \boxed{L_{\text{original}}} + \boxed{\text{reg}_d}$$

Original loss Regularization

Encourage Disentanglement: Loss Function

$$v_1 = [1, 2, 3, \dots]$$

$$v_2 = [5, 6, 7, \dots]$$

$$L_{\text{new}} = \boxed{L_{\text{original}}} + \boxed{\text{reg}_d}$$

Original loss Regularization

Can be any loss function
(e.g. BCE, W-Loss)

Encourage Disentanglement: Loss Function


$$\begin{array}{cc} & z_1 & z_2 \\ v_1 = [& 1, & 2, & 3, \dots] \\ v_2 = [& 5, & 6, & 7, \dots] \end{array}$$

$$L_{\text{new}} = \boxed{L_{\text{original}}} + \boxed{\text{reg}_d}$$

Original loss Regularization

Can be any loss function
(e.g. BCE, W-Loss)

Encourage Disentanglement: Loss Function



$$L_{\text{new}} = \boxed{L_{\text{original}}} + \boxed{\text{reg}_d}$$

Original loss Regularization

Can be any loss function
(e.g. BCE, W-Loss)

Summary

- Disentangled Z-spaces let you control individual features by corresponding z values directly to them
- There are supervised and unsupervised methods to achieve disentanglement

